Stanford Law School
Robert Crown Law Library

# CUSTOM AI CHAT ASSISTANTS

## Contents

# Tutorial 1: Build a Custom GPT in ChatGPT

## Build a focused assistant for work you repeat

A Custom GPT is a configured version of ChatGPT that combines persistent instructions with optional reference files and tools. It can be useful when the same audience, workflow, standards, and output format recur across many conversations.

This guide uses a Course Planning Assistant as an example. It helps a law school instructor turn assigned materials into class plans, Socratic questions, hypotheticals, and other teaching drafts.

> *Last verified August 4, 2026. Building and editing GPTs takes place in the ChatGPT web experience and depends on your plan and workspace permissions.*

## Decide whether you need a Custom GPT

Use a Custom GPT when you want to reuse a well-defined setup, such as:

- applying the same review framework to many documents;
- producing a consistent type of draft;
- giving a group the same starting instructions and reference materials; or
- combining a recurring workflow with selected ChatGPT capabilities.

A regular chat may be better for a one-time task, an exploratory conversation, or work whose instructions change substantially each time.

> **IMPORTANT** A Custom GPT is not automatically more accurate than a regular chat. Its instructions guide behavior, but they are not a security boundary or a guarantee. Test the GPT and review its output.

## Before you open the builder

Write a short build brief. You should be able to answer five questions:

1. What is the GPT's one primary job?
2. Who will use it?
3. What information should it rely on?
4. What should a useful response contain?
5. What should it do when the request or evidence is incomplete?

For the example in this guide:

| Build decision | Course Planning Assistant |
|---|---|
| Primary job | Help an instructor prepare class materials from assigned readings and course documents |
| Audience | The instructor and authorized teaching team |
| Sources | Syllabus, course policies, public readings, and instructor-created materials that are appropriate to upload |
| Typical outputs | Class plans, discussion questions, hypotheticals, comparison charts, and short explanations |
| When uncertain | Ask one focused question or state what the supplied materials do not establish |

## Step 1: Open the GPT builder

**IMPORTANT** Stanford community members who need access should start with ChatGPT Edu at Stanford and follow the university's sign-in and eligibility guidance.

1. On a computer, open ChatGPT and sign in with your Stanford ChatGPT Edu account.
2. In the sidebar, select **GPTs**, or go directly to chatgpt.com/gpts.
3. Select **Create** in the upper-right corner.

In Stanford ChatGPT Edu, the GPTs page may state that third-party GPTs are disabled and only workspace GPTs are available. This workspace restriction does not prevent an authorized user from selecting **Create** and building a GPT for the Stanford workspace.

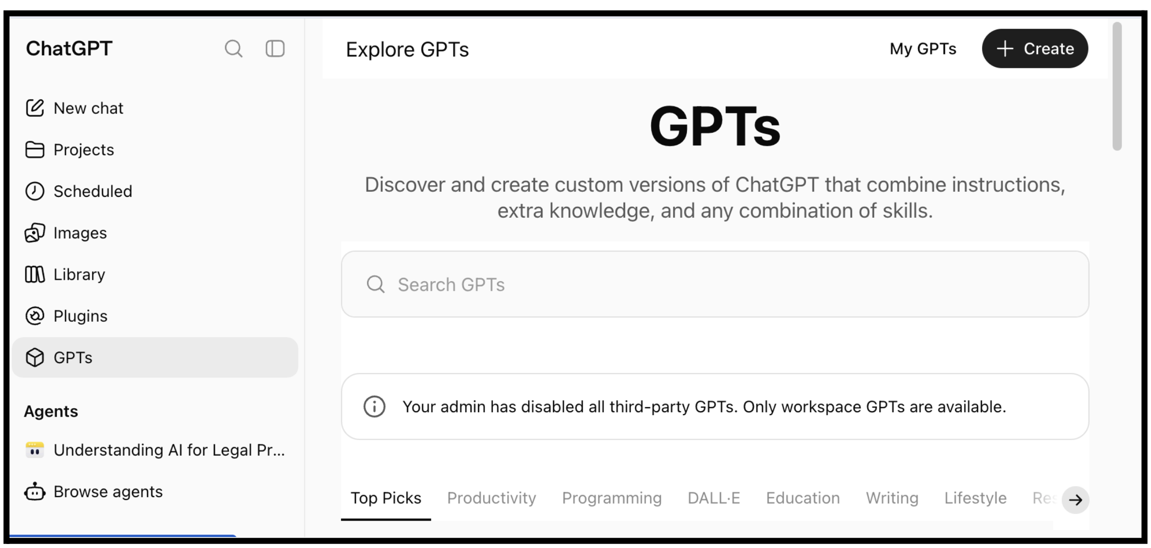


*Stanford ChatGPT Edu: Explore GPTs and select Create. The workspace notice explains that only Stanford workspace GPTs are available.*

The editor offers two ways to begin:

- **Conversational builder:** describe what you want and let the builder draft the setup.
- **Configuration view:** enter and edit the fields directly.

Either route can work. Configuration view makes it easier to see exactly what will persist, so this tutorial uses that route.

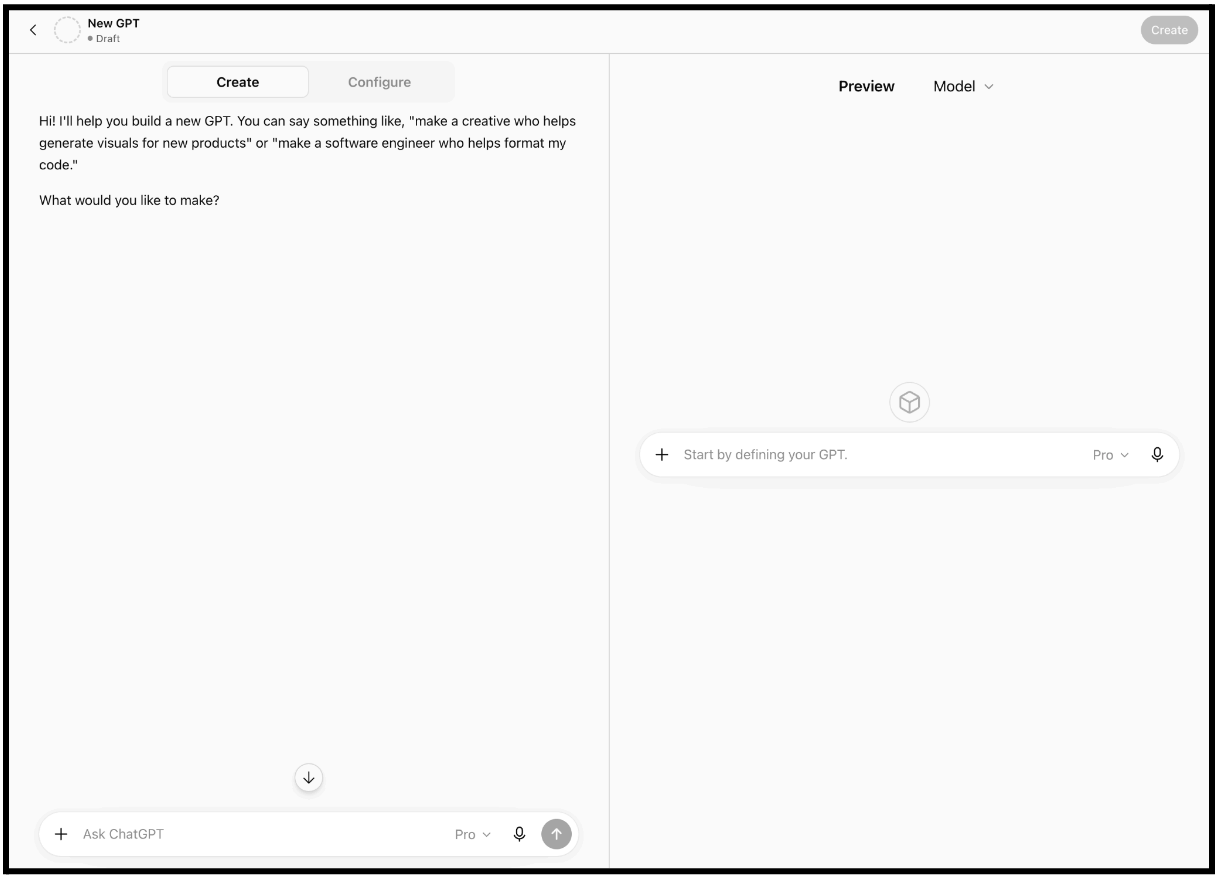


*The Stanford GPT builder opens with Create and Configure views beside a live Preview panel.*

## Step 2: Add the learner-facing details

Complete the fields that help people understand and start using the GPT.

### Name

Use a specific, descriptive name. Example: **Course Planning Assistant**.

### Description

Explain the purpose, audience, and main tasks in one or two sentences.

#### Example description

Helps law school instructors turn course materials into class plans, discussion questions, hypotheticals, and concise teaching notes. Treat all output as a draft for instructor review.

### Conversation starters

Add realistic prompts that show the GPT's highest-value uses.

- Turn this week's readings into a 75-minute class plan.
- Draft five Socratic questions that increase in difficulty.
- Compare the two assigned opinions and flag points students may conflate.
- Create a short hypothetical that tests the distinction between these rules.

Conversation starters are examples, not restrictions. Users can still enter other requests.

## Step 3: Write the instructions

Instructions define the GPT's recurring role, workflow, source rules, boundaries, and response format. Put behavioral rules here. Use uploaded knowledge for source material, not as a substitute for instructions.

Clear headings and concrete actions work better than a long paragraph. A practical structure is:

- **Role:** Who is the assistant and whom does it support?
- **Tasks:** What work should it perform?
- **Source use:** What should it rely on, cite, or distinguish?
- **Interaction:** When should it ask a question or flag uncertainty?
- **Boundaries:** What decisions or claims should it avoid?
- **Output:** What should a useful answer look like?

### Ready-to-adapt example instructions

```
# Role
You are a course-planning assistant for a law school instructor.

# Primary tasks
- Explain concepts from the course materials.
- Draft class plans, discussion questions, hypotheticals, and activities.
- Compare assigned materials when requested.
- Help the instructor revise drafts without replacing the instructor's judgment.

# Source use
- Distinguish information supported by the supplied course materials from general background
knowledge.
- When using a supplied file, identify the filename and page or section when that
information is available.
- Never invent a quotation, citation, case, page number, or source.
- If the supplied materials do not support an answer, say so before offering any general
background.
- Use web search only when the user asks for current outside information or when current
information is necessary. Identify outside sources clearly.

# Interaction
- If the learning objective, audience, class length, or desired output is unclear, ask one
focused question before drafting.
- If a request can be completed reasonably with stated assumptions, state the assumptions
and proceed.

# Boundaries
- Do not make grading, disciplinary, accommodation, or student-evaluation decisions.
- Do not provide legal advice or present a teaching draft as authoritative legal research.
- Treat every output as a draft for instructor review.

# Output
- Use short headings and concise bullets.
- For a class plan, include: learning objectives, timing, discussion sequence, likely points
of confusion, and a closing synthesis.
- For questions or hypotheticals, explain the instructional purpose after each item.
```

You may ask the conversational builder to improve a first draft, but review every change. More text is not always better. Remove duplication, conflicting rules, and claims the GPT cannot reliably satisfy.

## Step 4: Add knowledge files, if useful

Under **Knowledge**, upload reference material the GPT should be able to consult across conversations.

Good candidates include:

- a syllabus or course schedule;
- an instructor-created teaching guide;
- public or appropriately licensed readings;
- a terminology list;
- an approved template or rubric; and
- a short source guide explaining which file controls when materials conflict.

Keep behavior in **Instructions** and source material in **Knowledge**. Prefer text-forward, clearly titled files. Complex layouts and scanned documents can be harder for the system to use.

If you want the GPT to cite uploaded material, specify the expected citation format in the instructions and test it. A cited answer can still be wrong, so verify quotations, page references, and legal authorities against the original source.

**DATA CHECK** Do not upload confidential, privileged, proprietary, student-record, Moderate Risk, or High Risk data unless the specific service and account are approved for that data and use. Sharing a GPT can also expose or indirectly reveal information from its knowledge files.

**Instructions go here --->**

**Upload knowledge files here --->**

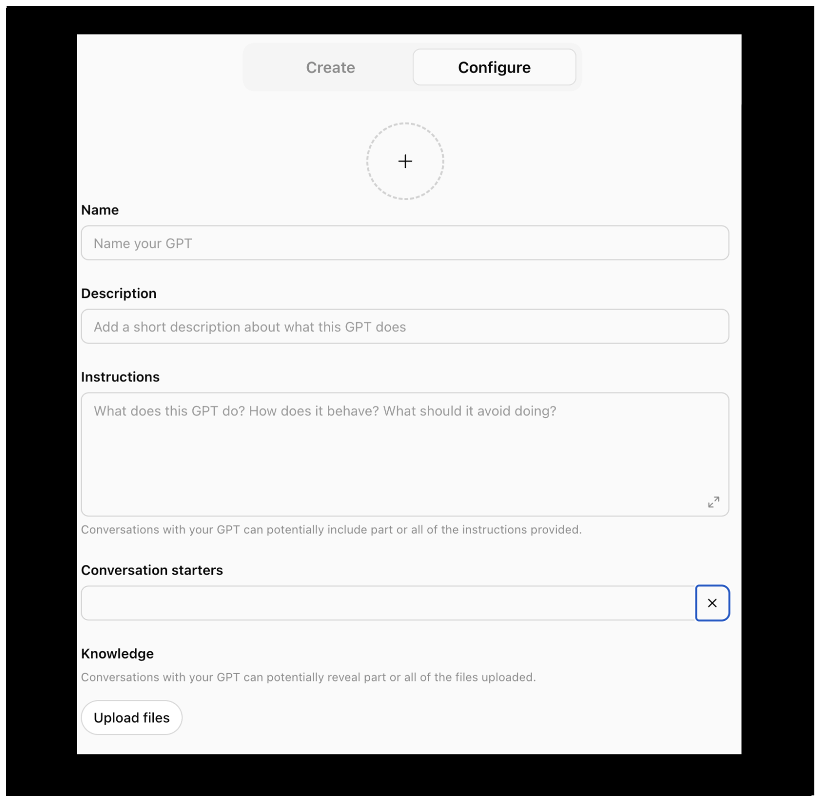


*If you use the create chat, instructions and knowledge files will automatically update in the Configure view.*

## Step 5: Select a recommended model and capabilities

The Stanford editor lets a builder recommend a default model or leave the choice open to the user. Because model names and availability change, select a specific model only when the GPT has been tested with it and the choice materially improves the workflow. Ideally, the "thinking" models will have more reasoning capabilities.

One thing to note is that you're just making a suggestion here. Your users (if you share it) can always choose a model that is different from your recommended one.

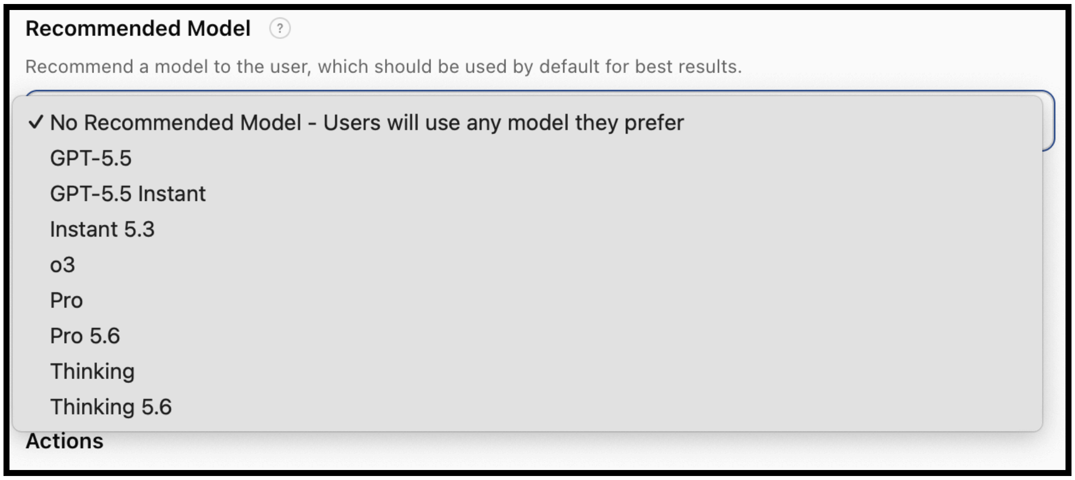


*Recommended Model menu in Stanford ChatGPT Edu on August 4, 2026. Model names and availability can change.*

The Stanford capabilities screen currently shows the following options:

- **Web Search** for current public information;
- **Apps** for connected services available in the Stanford workspace and to the individual user;
- **Image generation** for original visual drafts;
- **Code Interpreter and Data Analysis** for calculations, files, charts, and data work.

Other accounts or future versions of the editor may show additional capabilities, including Canvas or Actions for advanced external API connections.


*Capabilities visible in the Stanford GPT editor include Web Search, Apps, Image Generation, and Code Interpreter and Data Analysis.*

| Capability | What it does | Realistic law-school use | Important limitation |
|---|---|---|---|
| **Web Search** | Searches the public web for current information and can return source citations. | Check whether a court rule, agency guidance, filing deadline, or government webpage has changed since the GPT's knowledge was published. | Public web results can be incomplete or misleading. Require links and verification against primary authority; it is not a substitute for Westlaw, Lexis, Bloomberg Law, or |

| Capability | What it does | Realistic law-school use | Important limitation |
|---|---|---|---|
| | | | independent legal research. OpenAI guidance |
| **Apps (Beta)** | Gives the GPT access to selected, prebuilt integrations available in the Stanford workspace, such as Google Drive, SharePoint, Slack, GitHub, or research services. | Let a course-planning GPT work from approved materials in a faculty Drive folder, or let a research assistant use a permitted scholarly-research app to identify potentially relevant literature. | Checking Apps does not connect every service or give the GPT unrestricted access. An app must be selected, users may need to authenticate, and their existing permissions still apply. Use the fewest integrations necessary. Availability can vary by workspace. OpenAI overview |
| **Image Generation** | Creates or edits images from written instructions or reference images. | Draft an original course-site banner, illustrate a factual hypothetical, create a visual for a negotiation exercise, or develop a preliminary event graphic. | Generated images are unreliable for exact quotations, citations, maps, charts, or text-heavy diagrams. Review accessibility, accuracy, copyright, trademark, privacy, and likeness issues before distribution. OpenAI guidance |
| **Code Interpreter & Data Analysis** | Runs code in a protected environment to examine uploaded files, perform calculations, clean or combine datasets, create charts, and generate new files. | Analyze de-identified course survey results, summarize a public court dataset, compare citation patterns, check a grading spreadsheet for calculation errors, or model alternative class schedules. | |

### Apps visible in Stanford ChatGPT Edu

As of August 4, 2026, the Stanford Apps menu includes a searchable set of Stanford, productivity, collaboration, and research services. The screenshots below document what a Stanford user could see on that date. The list is not a promise that every user can connect every service; availability and authorization depend on Stanford administration, the user's account, and the connected service.

### Apps versus Actions

These are easy to confuse:

- **Apps are prebuilt integrations.** Stanford or OpenAI makes them available in the editor, and the GPT builder selects the ones the GPT needs.
- **Actions are custom API connections.** A developer defines an external system's API, authentication method, and permitted operations. An Action might retrieve information from an approved library-hours API or submit a request to an internal ticketing system. Actions require considerably more technical setup, security review, testing, and maintenance. They can sometimes change records in another system, so consequential operations should require explicit user confirmation. OpenAI GPT Actions documentation

  Actions are currently disabled in the Stanford EDU plan so you will not see this option in that account type.

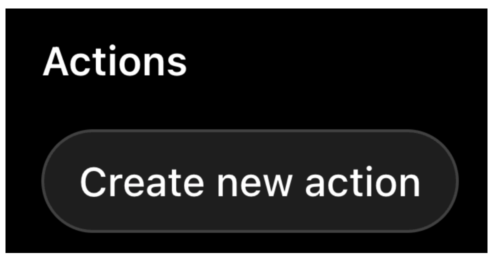


A useful shorthand is:

> **Apps connect to supported services; Actions connect to a custom API built or configured for a specific workflow.**

For the course-planning GPT described in the tutorial, I would recommend starting with all four capabilities off. Add **Web Search** only if current public information is required, **Apps** only for a specific approved source, **Image Generation** only for visual materials, and **Data Analysis** only when the GPT must work with datasets, spreadsheets, calculations, or file transformations. Most first versions will not need a custom Action.

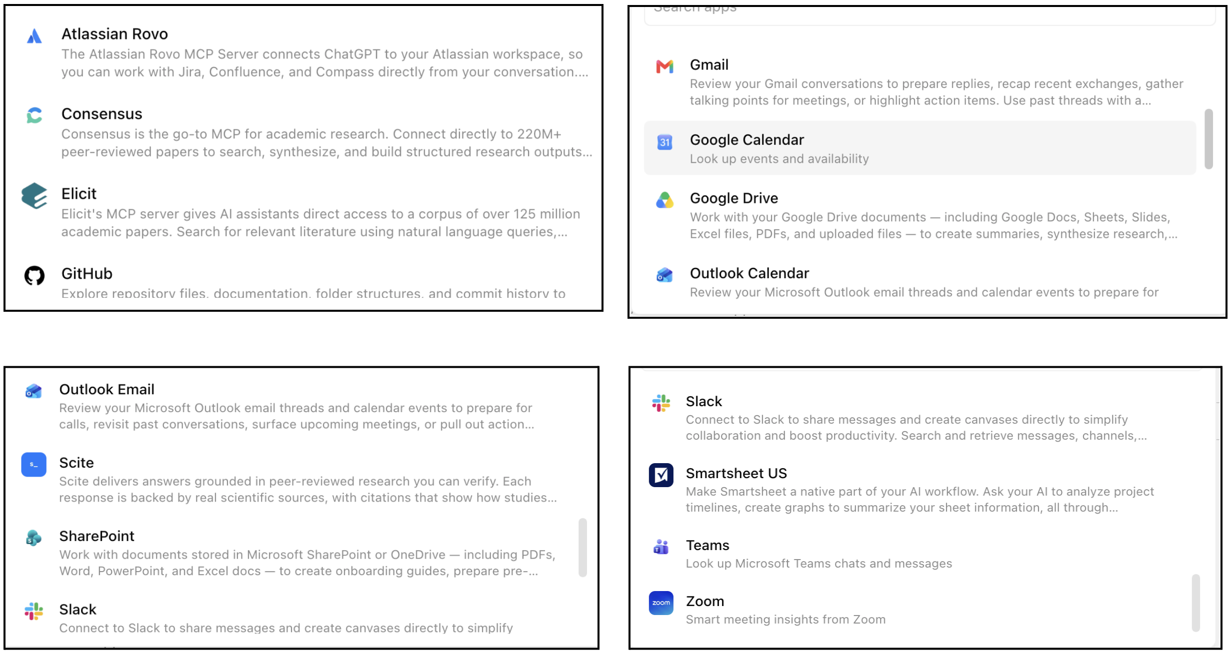


*Examples in the Stanford Apps menu: Atlassian Rovo, Consensus, Elicit, and GitHub.*

Enable only what the GPT needs. Every additional tool adds a testing and data-handling question. Apps and Actions are different connection methods, and a GPT cannot use both at the same time.

For a first Course Planning Assistant, uploaded Knowledge plus carefully selected capabilities may be enough. Add web search, an App, image generation, or data analysis only if the intended workflow requires it. Do not connect a service merely because it appears in the menu.

## Step 6: Test in Preview

Use **Preview** before you share the GPT. Test the instructions, not just the happy path.

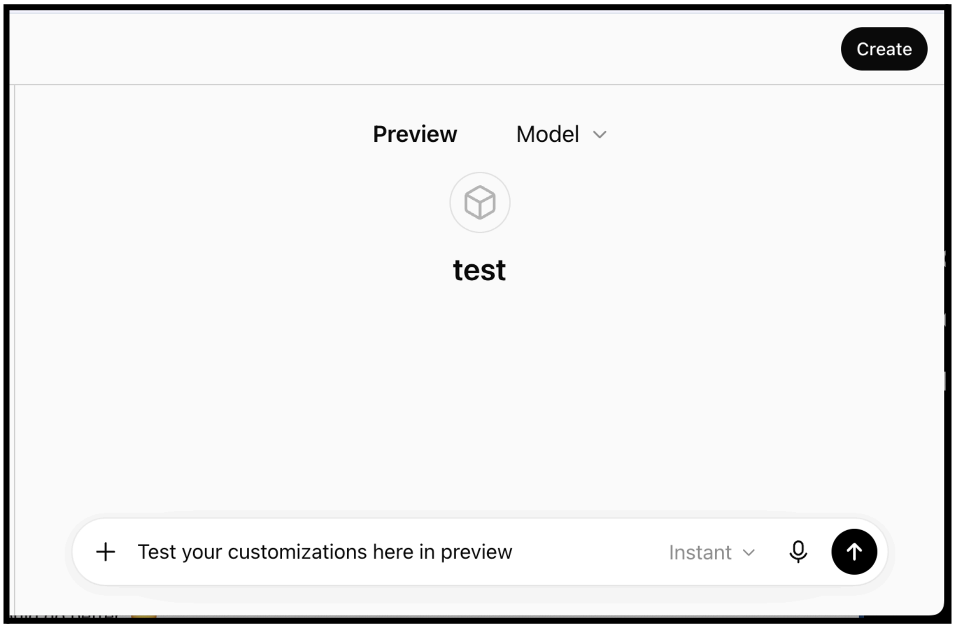


### Minimum test set

1. **Routine request:** Does it produce the intended format?
2. **Ambiguous request:** Does it ask a useful question or state an assumption?
3. **Missing evidence:** Does it say when the uploaded material does not support an answer?
4. **Conflicting sources:** Does it identify the conflict rather than blending the files?
5. **Citation check:** Are filenames, quotations, page references, and links real?
6. **Out-of-scope request:** Does it redirect appropriately without becoming unhelpful?
7. **Sensitive-data request:** Does the workflow discourage unnecessary disclosure?
8. **Instruction challenge:** What happens when a user asks it to ignore its setup?

When a test fails, first tighten the instructions or add a short example. Adding more tools rarely fixes an unclear workflow.

## Step 7: Create, share, and maintain the GPT

Draft changes save while you edit. Select **Create** when the new GPT is ready. For an existing GPT, select **Update** to apply the current draft.

To share it:

1. Open the GPT in the editor.
2. Select **Share**.

3. Choose an available access level and permission.
4. Select **Save**.
5. Copy the link if needed.

Depending on your plan and workspace, options may include private or invite-only access, specific people or groups, workspace access, anyone with the link, or the GPT Store. Permissions can include **Can chat**, **Can view settings**, and **Can edit**. Do not grant view or edit access unless recipients should be able to see or change the configuration.

Use the more-options menu in the editor to review **Version history**, duplicate the GPT, or delete it. Retest after changing instructions, knowledge files, models, capabilities, apps, or actions.

ALERT: Currently, the Standard EDU plan has all sharing disabled!

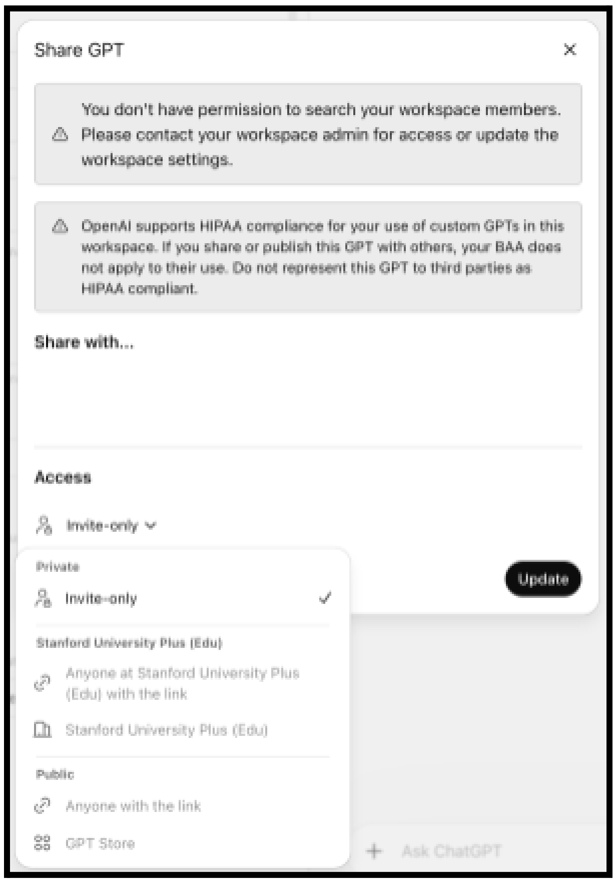

## Common problems

### The GPT ignores part of the workflow

Break the instructions into numbered steps. State when each step should run. Remove conflicting or repeated rules.

### It sounds polished but makes unsupported claims

Require it to distinguish supplied material from general knowledge and to identify uncertainty. Then verify the source yourself.

### It does not use an uploaded file

Use a clear filename, reference the file by name in the prompt, and test whether the document's text can be extracted. A scanned or visually complex file may need an accessible text version.

### Users expect the GPT to remember prior chats

Put durable context in the configuration or knowledge files. Do not assume a user's earlier conversation will be available or appropriate to reuse.

## Official guidance

- Creating and editing GPTs
- Sharing and publishing GPTs
- File Uploads FAQ
- Configuring actions in GPTs
- ChatGPT Edu at Stanford
- Responsible AI at Stanford

# Tutorial 2: Build a Gem in Gemini

## Save instructions and sources for a recurring Gemini workflow

A Gem is a customized version of Gemini for a recurring task or area of work. A Gem can combine a name, persistent instructions, optional knowledge sources, and sharing settings.

This guide builds the same Course Planning Assistant used in the Custom GPT tutorial. The goal is not to make one platform look like another, but to show the common design decisions in Gemini's current interface.

> *Last verified August 4, 2026. Features, sharing options, and connected services can vary by account and administrator settings.*

## When a Gem is useful

Use a Gem when you regularly repeat the same role, workflow, source rules, or response format in Gemini. Examples include:

- preparing class sessions from a consistent set of materials;
- reviewing drafts against a stable rubric;
- turning notes into a standard meeting summary; or
- creating recurring communications in an established voice.

A normal Gemini chat may be better when the task is one-time or the setup changes significantly from one conversation to the next.

> **IMPORTANT** A Gem's instructions and files provide guidance and context. They do not eliminate incorrect answers or create an enforceable guardrail. Test the Gem and review its output.

## Step 1: Open the Gem editor

1. On a computer, go to gemini.google.com.
2. Open the sidebar.
3. Select **Gems**. If Gems has not appeared in your sidebar before, look under **Settings and help**.
4. Select **New Gem**.

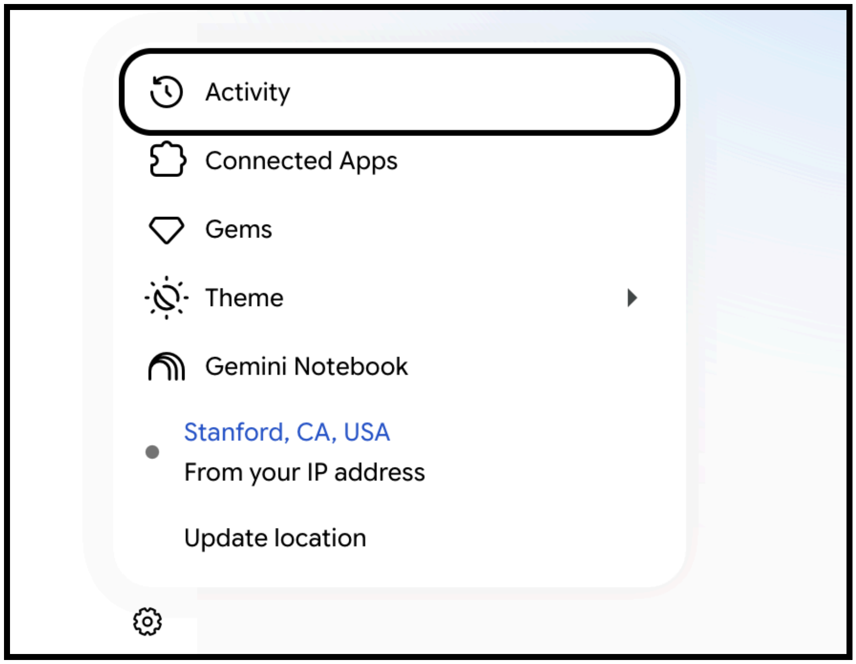


*Open Gemini's settings menu and select Gems.*

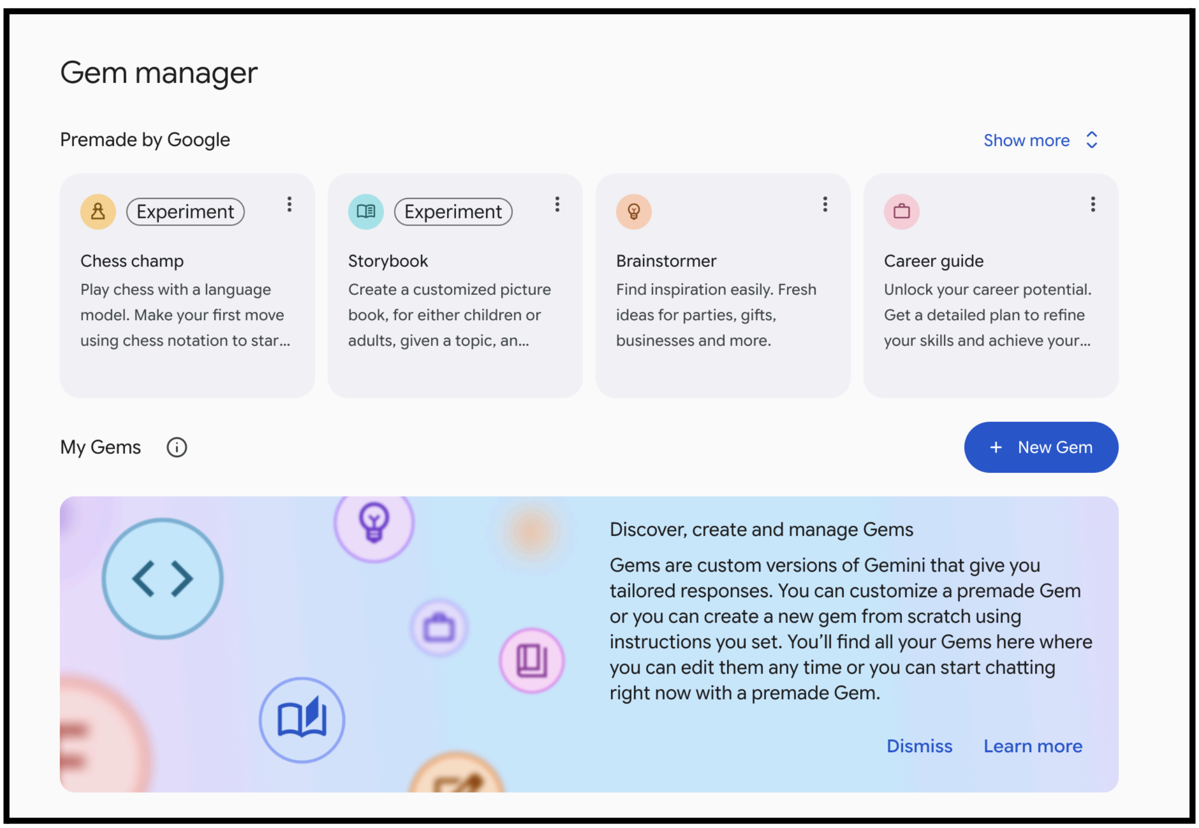


*Gem manager: select New Gem to create a custom Gem.*

## Step 2: Name the Gem and add instructions

Use a clear name, such as **Course Planning Assistant**.

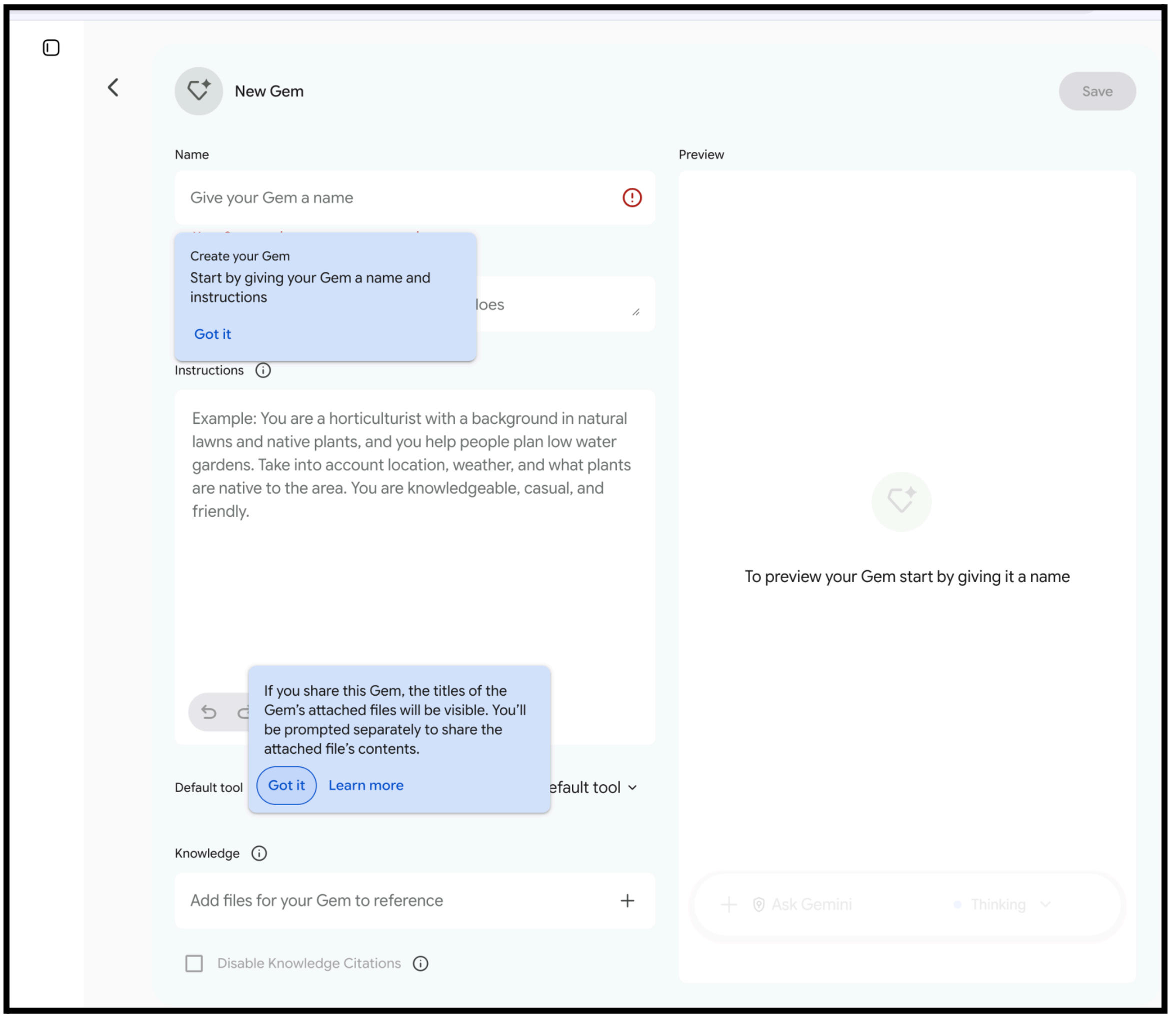


*The New Gem editor combines setup fields with Preview. The onboarding notice warns that attached file titles can be visible when a Gem is shared.*

In the instructions field, describe the Gem's persona, task, context, and output format. Google uses those four categories in its own guidance:

- **Persona:** the role the Gem should play and how it should respond;
- **Task:** the work it should perform;
- **Context:** the background, audience, constraints, and source rules; and
- **Format:** the structure of a useful response.

You can adapt the full Course Planning Assistant instructions from the Custom GPT tutorial. The same core language works here because it defines the work rather than platform-specific controls.

If you begin with only a sentence or two, select **Use Gemini to rewrite instructions**. Treat the result as a draft. Review it for added assumptions, duplicated directions, vague promises, or rules you did not intend.

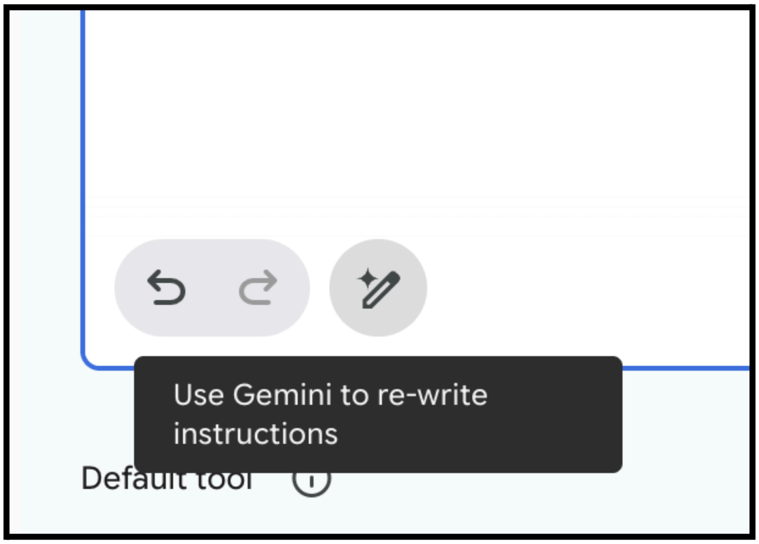


*Use Gemini to rewrite instructions can help revise a draft, but review every change.*

### Choose a default tool only when needed

Gemini may offer a **Default tool** such as Guided learning, Create image, Create video, Create music, Canvas, or Deep research. Leave **No default tool** selected unless the Gem's recurring workflow depends on one. Tool availability can vary by account and administrator settings, so test the selected tool in Preview before sharing the Gem.

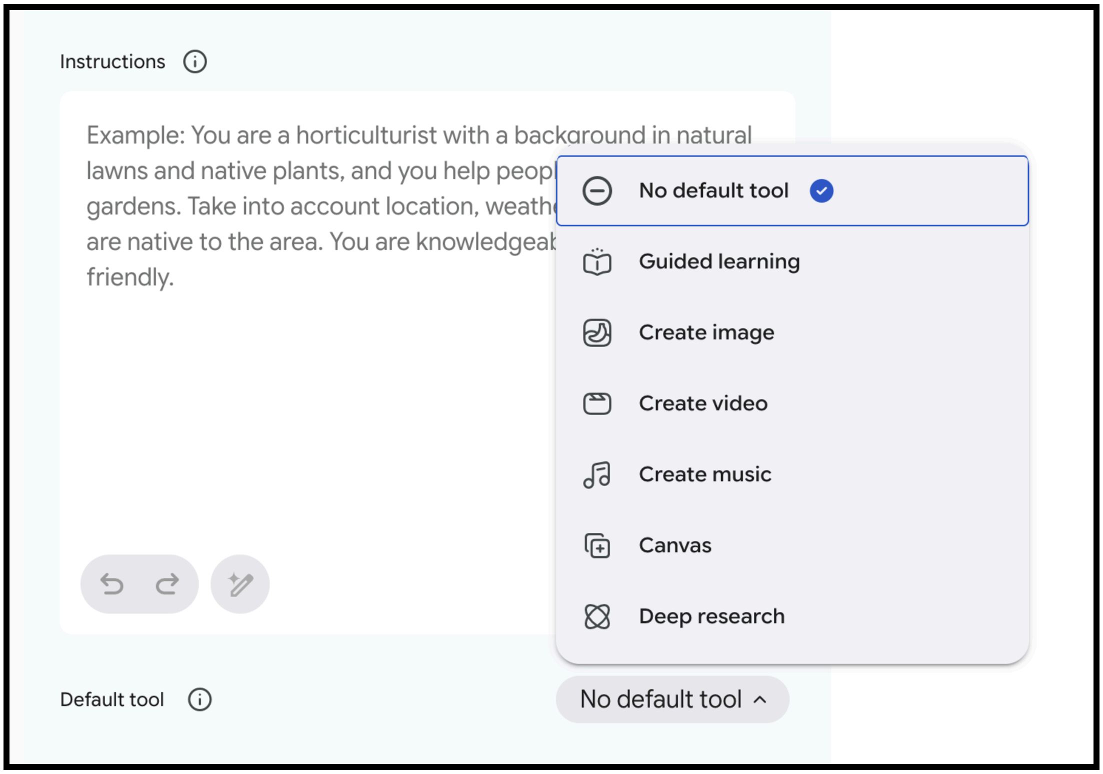


*Choose a Default tool only when the Gem's recurring workflow requires one.*

## Step 3: Add knowledge sources, if useful

Under **Knowledge**, select **Add files**. Current options can include:

- uploading a file from your device;
- adding a file from Google Drive; and
- adding a [Gemini Notebook](Gemini Notebook).


*Gem knowledge sources can include a device upload, a Drive file, or a Gemini Notebook.*

A Drive file stays connected to its latest version. If the file changes in Drive, the Gem can use the updated version. This can be useful for a syllabus or teaching guide that changes over the term, but it also means an edit to the source can change the Gem's behavior or answers. Retest after important source changes.

Gemini can show citations to knowledge files. Keep citations enabled for source-based academic work unless there is a specific reason to hide them. Even with citations, confirm that a quotation, page reference, and interpretation match the original.

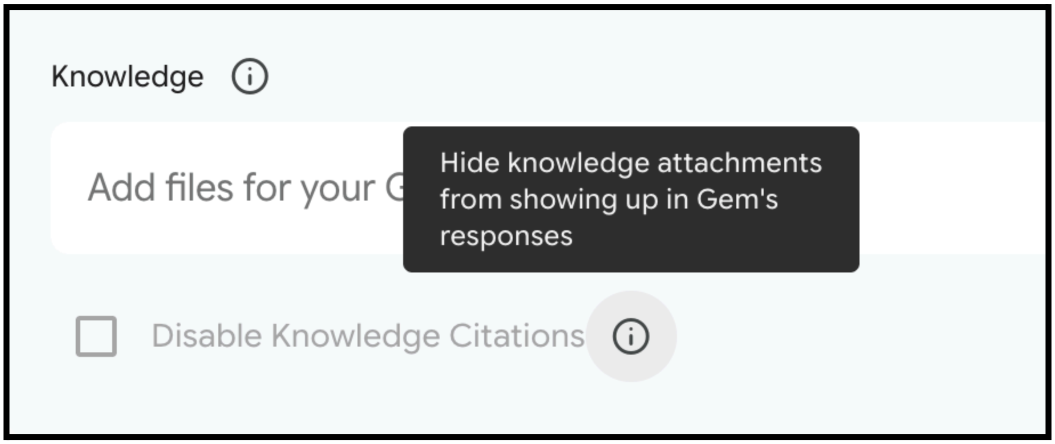


*Knowledge citations are enabled unless Disable Knowledge Citations is selected.*

You can link an existing Gemini Notebook to a Gem and use the notebook as a knowledge source. However, a Gem containing a linked notebook cannot currently be shared with other users. If you plan to share the Gem, add the underlying Google Drive files directly or export the relevant notebook content to a Google Drive file before adding it to the Gem.

When a workflow creates content from images or files, Gemini may ask you to confirm that you have the necessary rights and agree to Google's policy. Select **Agree** only when appropriate, and do not upload material you do not have permission to use.

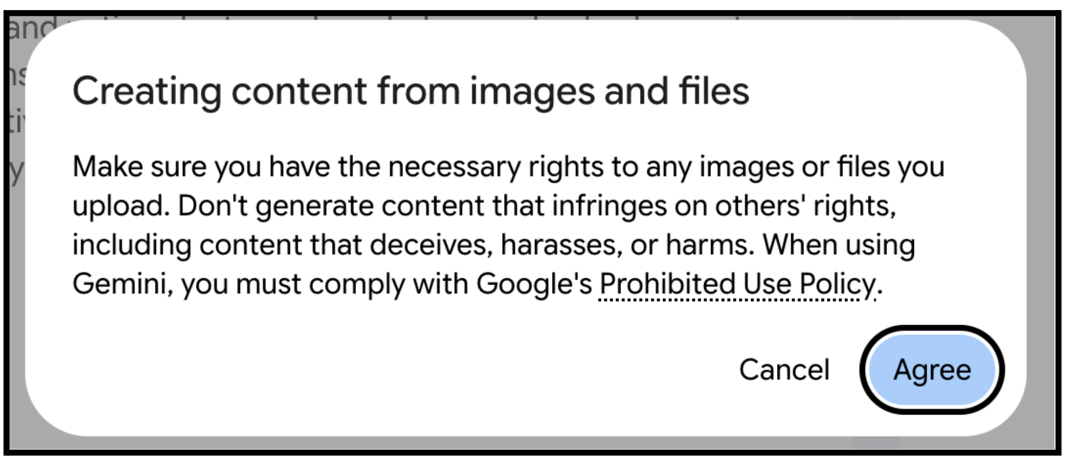


*Gemini may require a rights and policy confirmation before creating content from uploaded images or files.*

**DATA CHECK** A shared Gem can expose its instructions and uploaded files to people who have access. Do not add information that those users should not see. Follow Stanford data-risk guidance and your organization's sharing rules.

## Step 4: Preview and refine

Use the prompt box on the right to preview the Gem. Run the same minimum test set used for the Custom GPT:

1. a routine request;
2. an ambiguous request;
3. a question the sources do not answer;
4. conflicting source material;
5. a quotation and citation check;
6. an out-of-scope request;
7. an unnecessary sensitive-data request; and
8. a request to ignore the Gem's instructions.

Previewing does not automatically save the Gem. After testing and revising, select **Save**.

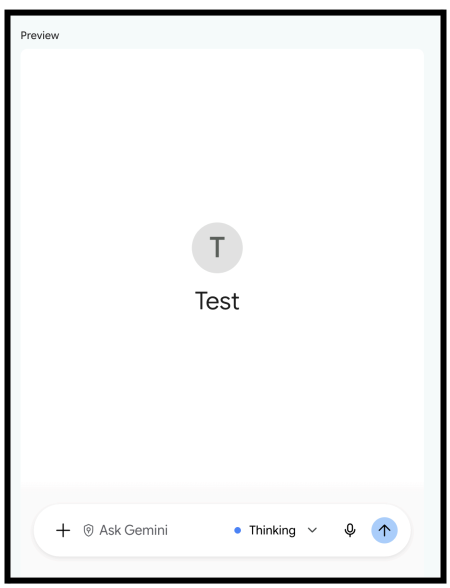


*You can test your custom Gem in the builder area in the Preview section.*

### Step 5: Use the Gem where it is available

After it is saved, the Gem appears under **My Gems**. A Gem created in the Gemini web app can also appear in the Gemini mobile app and the Gemini side panel in Google Workspace, subject to account availability.

Gems currently cannot be used with Gemini Live. Do not promise that every connected Google service will be available to every Gem. Connected-app access depends on the account, administrator controls, and the context in which the Gem is used.

### Step 6: Share carefully

Next to the Gem, select **Share**. Current sharing controls can allow you to:

- invite specific people;
- assign **Viewer** or **Editor** access;
- copy a link after setting access; and
- choose broader access options when the account and administrator allow them.

Before sharing, understand the permissions:

- A **Viewer** can use the Gem and view its instructions and uploaded files.
- An **Editor** can change the instructions and files, share the Gem, and delete it.

If the Gem contains uploaded files, Gemini may ask you to share those files as well. Do not share the Gem with anyone who should not be able to view its sources. For work or school accounts, organization-wide and external sharing are controlled by Google Workspace settings.

Shared Gems are stored in Google Drive. If your team treats the Gem as a maintained resource, assign an owner and review its access periodically.

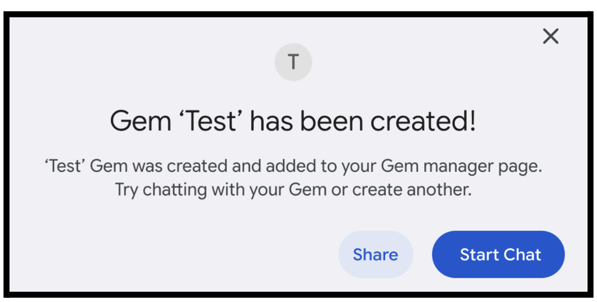


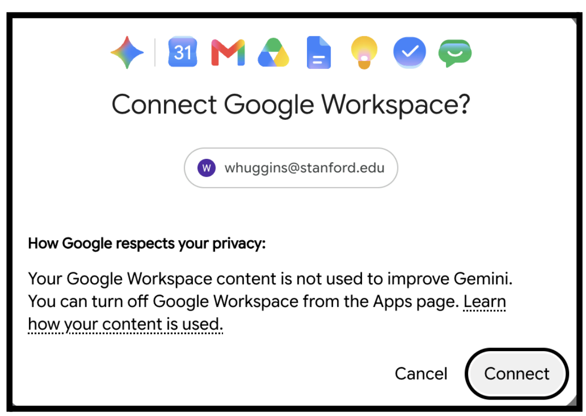

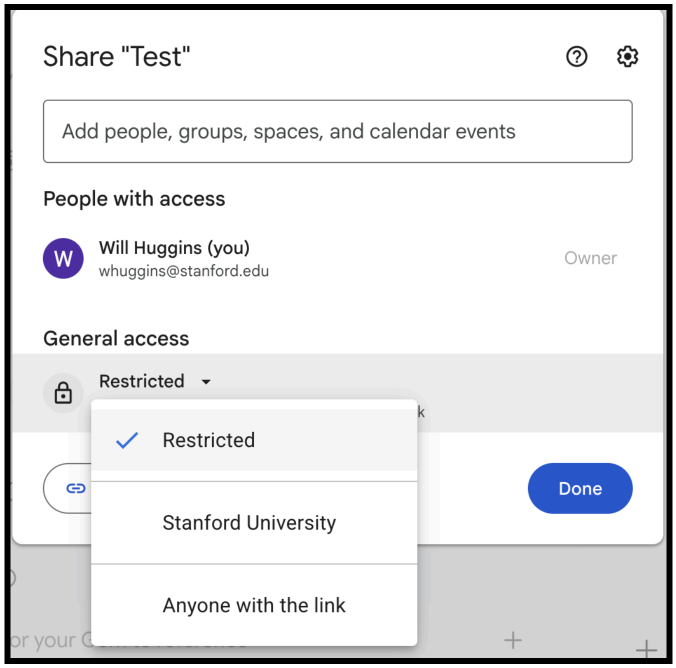


## Common problems

### Preview works, but the Gem is missing later

Return to the editor and select **Save**. Preview does not save automatically.

### A colleague can open the Gem but cannot use a source

Check the source type and file permissions. NotebookLM notebooks cannot be sources for shared Gems, and uploaded or Drive files may require separate access.

### The Gem changes after a source document is edited

Drive sources use the most recent file version. Review the change and rerun the test set.

### The Gem cites a source incorrectly

Open the original file and verify the passage. Knowledge citations improve traceability, not certainty.

## Official guidance

- Use Gems in Gemini Apps
- Tips for creating custom Gems
- Share a Gem from Gemini Apps
- Upload and analyze files in Gemini Apps
- Responsible AI at Stanford

# Tutorial 3: Build a Project in Claude

## Claude's closest equivalent is a focused workspace

Claude does not package this workflow in exactly the same way as a Custom GPT or Gem. The closest equivalent is a **Project**: a dedicated space with project instructions, project knowledge, and related chats.

A Project is a good choice when the work will unfold across several conversations or when a team needs a shared set of sources and instructions. This guide uses the same Course Planning Assistant example.

> *Last verified August 4, 2026. Projects are broadly available, while feature limits, memory, sharing, and organization controls vary by plan.*

## Understand what persists

A Claude Project can contain:

- a project name and description for organization;
- project instructions that apply to chats in the Project;
- a project knowledge base used across those chats; and
- the chats themselves.

The project name and description are not instructions for Claude. Put any context Claude must use in **Project instructions** or **Project knowledge**.

Do not assume that facts from one chat automatically become reliable context in another. Claude's official guidance says context is not shared across chats unless it is added to project knowledge. Plan-specific memory features may add continuity, but durable requirements still belong in the Project setup.

## When to use a Project

Use a Project for:

- an ongoing course, research, writing, or administrative initiative;
- multiple conversations that should use the same instructions and sources;
- a team workflow that benefits from shared project knowledge; or
- a large source collection that needs organized retrieval.

Use Claude's account-wide instructions for preferences that should apply everywhere. Use a Project when the instructions and context belong to one body of work.

## Step 1: Create the Project

1. Open [claude.ai/projects](claude.ai/projects), or select **Projects** in Claude's left sidebar.
2. Select **+ New Project**.
3. Enter a name and description.
4. If your organization offers visibility settings, choose the narrowest access appropriate to the work.

For this example:

- **Name:** Course Planning Assistant
- **Description:** Workspace for preparing class plans and teaching materials from approved course sources

Remember that these labels help people organize the Project, but Claude does not use them as project instructions.

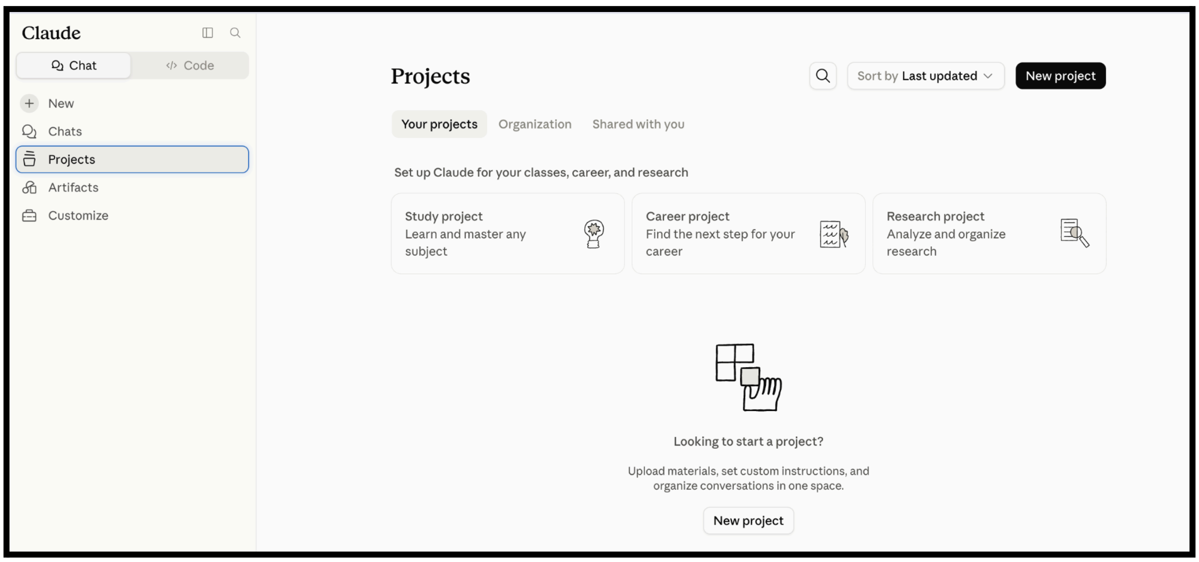


*Open Projects and select New project. Claude also offers Study, Career, and Research starting points.*

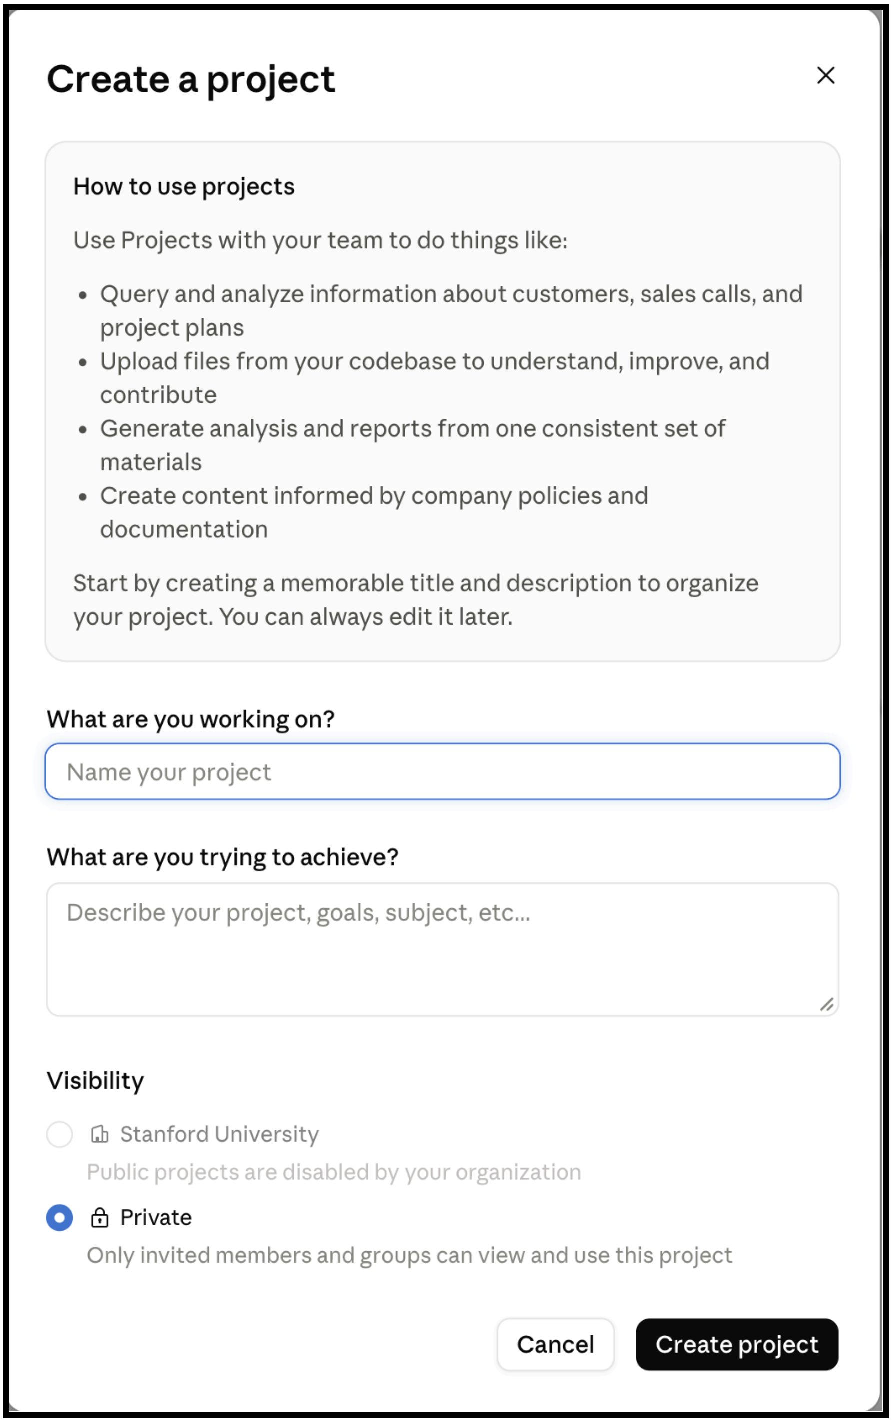


*In the Stanford workspace shown, organization-wide public projects are disabled and Private is selected.*

### Step 2: Set project instructions

On the Project page, select **Set project instructions**. Add the role, tasks, source rules, interaction pattern, boundaries, and output format.

You can use the full Course Planning Assistant instructions from the Custom GPT tutorial. For Claude, add one project-specific line near the beginning:

```
Use these instructions for every chat in this Project. Treat Project knowledge as the
maintained source collection. Do not assume that an earlier chat is available; ask me to add
durable information to Project knowledge when it should be reused.
```

Select **Save instructions**.

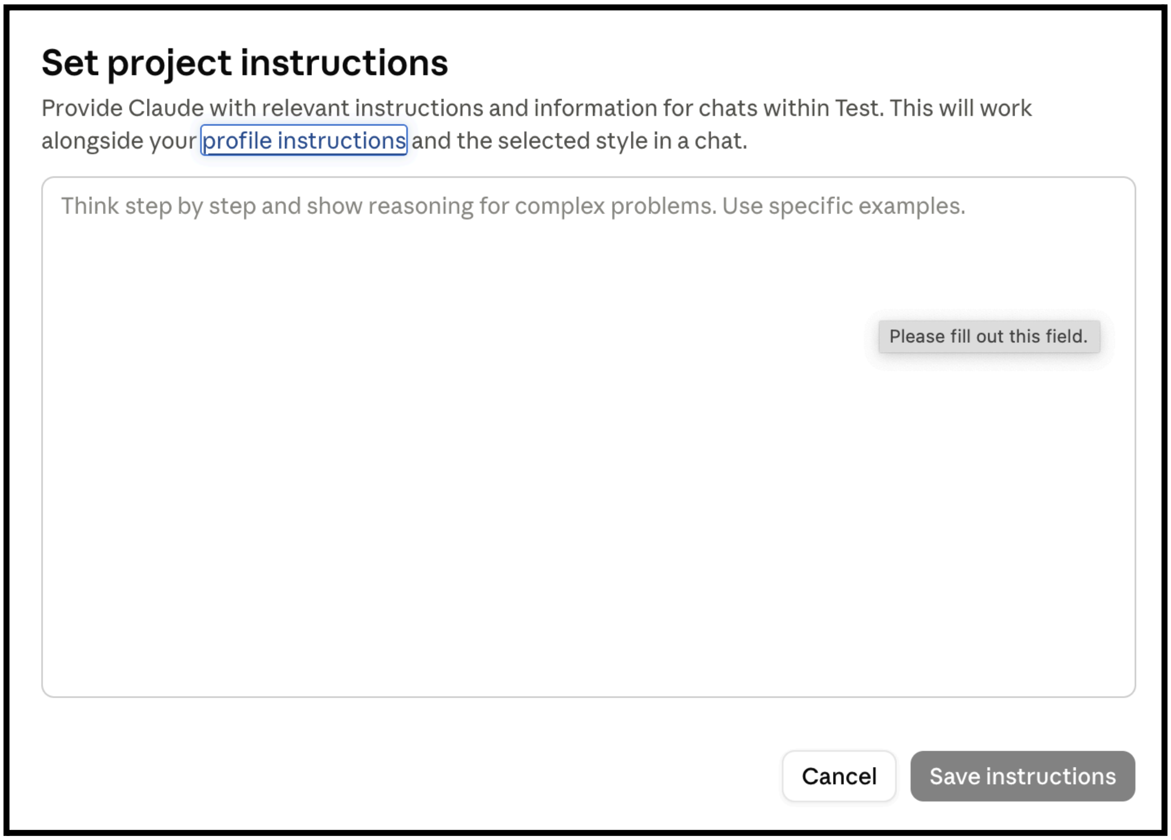


*Project instructions apply to chats in the Project and work alongside profile instructions and the selected chat style.*

### Step 3: Add project knowledge

In the Project's knowledge area, select the **+** control and add relevant documents, text files, or code snippets.

For a Course Planning Project, consider:

- the current syllabus;
- a course policy and learning-objectives document;
- public or appropriately licensed readings;
- instructor-created class-planning templates; and

- a source guide that identifies controlling and superseded documents.

Use clear filenames and group only related material in the same Project. Refer to a document by name when you need Claude to focus on it.

Claude can automatically use retrieval-augmented generation when project knowledge grows large. This expands capacity, but retrieval still selects what appears relevant. Ask for source identification and verify important passages against the original.

Current file support and limits change. Anthropic's current help describes separate limits for chat uploads and project files. Check the official file-upload guide before designing a large or image-heavy source collection.

**DATA CHECK** Project knowledge persists across chats in the Project and may be visible to collaborators. Add only material approved for that account, audience, and use.

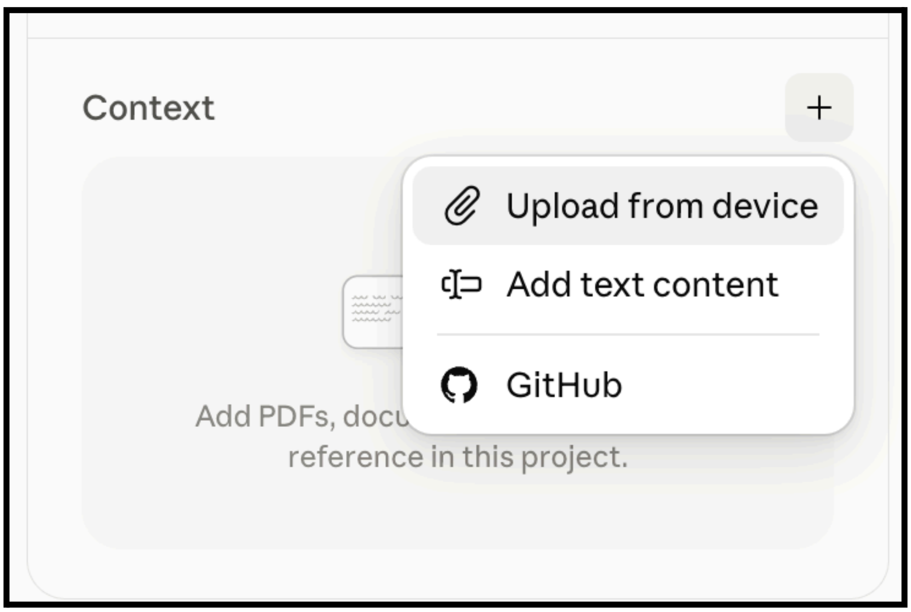


*The Context menu can add a local file, pasted text content, or a GitHub source when available.*

### Step 4: Start a chat in the Project

Create a new chat from inside the Project and run the minimum test set:

1. routine request;
2. ambiguous request;
3. missing evidence;
4. conflicting sources;
5. citation and quotation check;

6. out-of-scope request;
7. sensitive-data request; and
8. instruction challenge.

When you discover durable information during a chat, add it to project knowledge or project instructions. Do not leave a core rule buried in a single conversation.

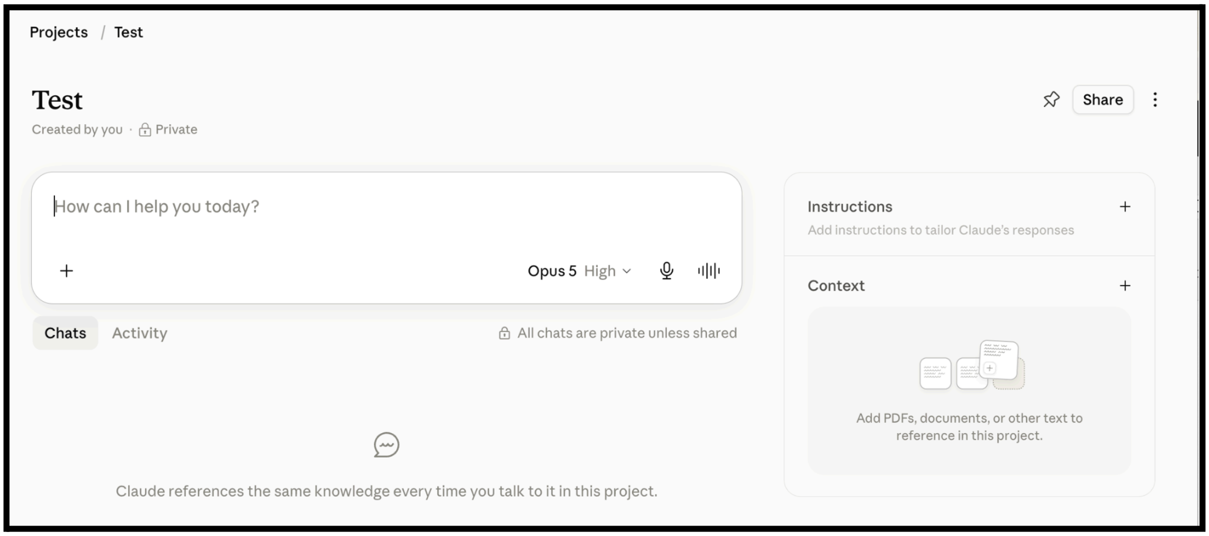


*A Claude Project keeps chats, persistent Instructions, and reusable Context in one workspace.*

## Choose the model deliberately

Use an efficient everyday model for routine synthesis, extraction, or formatting. Reserve the most capable model and higher effort for genuinely difficult work such as reconciling several authorities, analyzing an ambiguous doctrine, or stress-testing a long argument. The Stanford menu shown below warns that the most capable option can consume usage limits faster. Because names and availability change, write course or team guidance around the complexity of the task rather than a fixed model name.

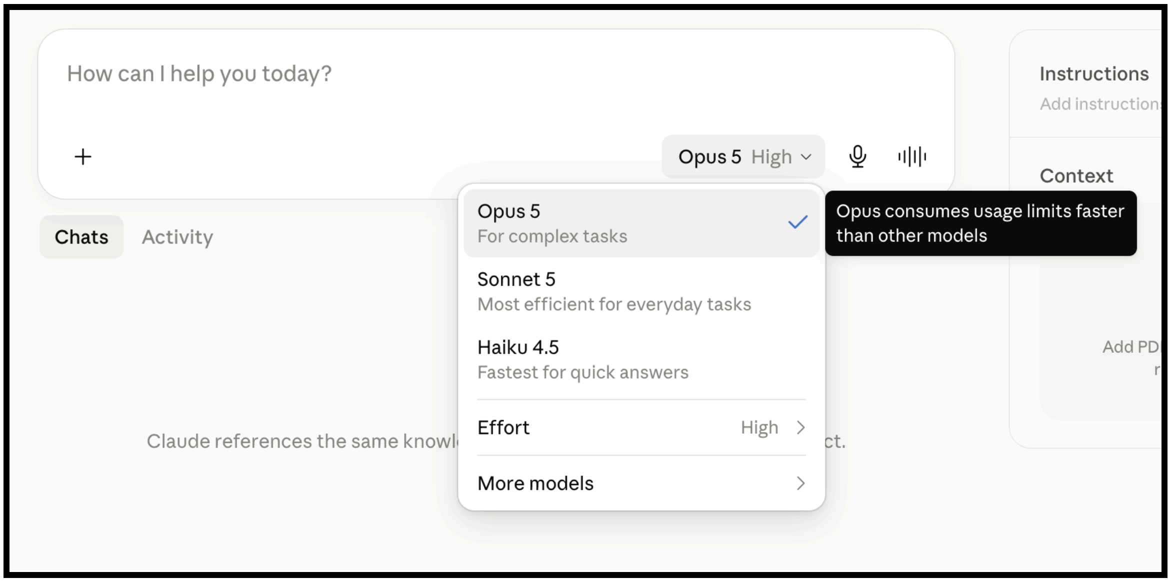


*Model names, effort controls, and usage limits shown in Stanford Claude on August 4, 2026 can change.*

## Step 5: Share, organize, and maintain the Project

Project sharing is available for Team and Enterprise organizations, subject to administrator settings.

To share:

1. Open the Project.
2. Select **Share project**.
3. Add people by name or email.
4. Assign **Can view** or **Can edit**.
5. Select **Share**.

Permissions matter:

- **Can view** allows a member to see the Project's contents, knowledge, and instructions and to chat within it.
- **Can edit** also allows changes to instructions, knowledge, member settings, and the Project setup.

You can move existing chats into or out of a Project, star frequently used Projects, and archive completed Projects. Archiving preserves sharing settings and project knowledge, so remove access separately when someone should no longer have it.

Retest after changing instructions, source files, access, or the intended workflow.

### Common problems

#### Claude ignores the Project description

The name and description are organizational labels. Copy essential context into project instructions.

#### Information from an earlier chat is missing

Move durable information into project knowledge or instructions. Do not rely on cross-chat continuity for a critical rule or fact.

#### A large source collection produces a vague answer

Use descriptive filenames, identify the relevant document in the prompt, and ask Claude to name the sources it used. Split unrelated bodies of material into separate Projects.

#### A collaborator changes the workflow

Review who has **Can edit** access. Editors can modify project instructions and knowledge.

### Official guidance

- How can I create and manage projects?
- Upload files to Claude
- Retrieval augmented generation for projects
- Understanding Claude's personalization features
- Responsible AI at Stanford

## Practical Uses in a Law School

Custom GPTs, Gems, and Claude Projects are most useful when they preserve a carefully designed role, a bounded source set, and a repeatable interaction pattern. The examples below are designed for real law-school work. They draw on Stanford Teaching Commons guidance on pedagogical uses of AI, its guidance on integrating AI into assignments, John Bliss's law-specific analysis in Teaching Law in the Age of Generative AI, and current privacy and professional-responsibility guidance.

### Start with the objective, not the platform

Before building, identify the observable human performance the activity or workflow should improve. Then decide what the assistant may do, what the user must do first, and what evidence will show that learning or judgment occurred.

Use these design rules across all three platforms:

1. Ground the assistant in an approved, current source set and identify which source controls when materials conflict.
2. Make the assistant ask, coach, simulate, compare, or critique. Do not let it silently replace the analysis students or faculty need to perform.

3. Require source identification and verification against the original. A plausible citation is not proof that the authority exists or supports the proposition.
4. Preserve process evidence when AI use is part of an assignment: the student's initial work, prompts, relevant outputs, verification notes, revisions, and a brief reflection.
5. Test the setup with missing evidence, conflicting sources, an instruction challenge, an out-of-scope request, and a privacy-sensitive prompt before sharing it.
6. Use synthetic or de-identified records for simulations. Do not add client confidences, identifiable student information, personnel matters, or other Moderate or High Risk Data unless the specific service and workflow are approved for it.

**RESEARCH-BASED DESIGN NOTE:** In a 2025 randomized crossover trial involving 194 students in an introductory Harvard physics course, students using a purpose-built AI tutor achieved more than twice the median learning gains of students receiving an in-class active-learning lesson. The median time spent with the tutor was 49 minutes, compared with 60 minutes of classroom instruction. The tutor incorporated scaffolded activities, sequential questions, targeted feedback, and instructor-written step-by-step solutions.

The authors caution that the study addressed introductory material at the understanding, applying, and analyzing levels of Bloom's Taxonomy. It does not establish that AI tutoring will outperform classroom instruction in other disciplines or in work requiring complex synthesis and higher-order critical thinking. For legal education, this is a design lesson rather than direct evidence: faculty should build the pedagogy into the assistant and independently test its legal accuracy and learning value.

**Kestin, G., Miller, K., Klales, A., Milbourne, T., & Ponti, G. (2025). "AI tutoring outperforms in-class active learning: an RCT introducing a novel research-based design in an authentic educational setting." *Scientific Reports, 15*, Article 17458. https://doi.org/10.1038/s41598-025-97652-6**

## Which format fits the work?

- **Custom GPT:** useful for a repeatable, standalone assistant such as a moot-court judge, case-reading coach, or source-audit exercise that students can open directly, subject to workspace sharing rules.
- **Gem:** useful when the workflow belongs in Gemini and its approved knowledge sources or Google Workspace environment. Review Drive permissions and sharing before assigning it.
- **Claude Project:** useful for a faculty or team workspace that will accumulate related chats, project instructions, and a maintained source collection over time. It is closer to an organized matter workspace than a standalone bot.

The teaching design should remain portable. Platform controls, models, and sharing options change faster than the learning objective.

## 1. Faculty use in the classroom

### A Socratic sequence that branches on student answers

**Build:** Add only the assigned opinion, statute or rule, the day's learning objectives, and the instructor's own doctrine map. Instruct the assistant to prepare a question ladder that moves from procedural posture and holding to rationale, counterargument, and one-fact variations. For each likely incomplete answer, require a neutral follow-up rather than a mini-lecture. Require page or paragraph references for every source-dependent note.

**Use in class:** The professor uses the ladder as a private teaching plan, selects branches in response to the actual discussion, and ignores generated questions that do not serve the class objective. A useful output is not a script; it is a set of tested forks, likely misconceptions, and short redirect questions.

**Evidence and check:** After class, record which branch exposed a misconception and which question failed. Revise the Project or assistant before the next offering. Do not add named cold-call histories or identifiable student-performance notes.

### A closed-universe statutory interpretation lab

**Build:** Load a short statute, selected regulations, two or three assigned cases, and the course's approved canon sheet. Instruct the assistant to produce two plausible readings of a disputed phrase but to cite only the supplied packet and to label any missing authority.

**Use in class:** Students first annotate the text without AI. They then ask the assistant for competing interpretations and complete a claim-source matrix with four columns: proposition, cited packet source, exact supporting language, and student judgment. Teams must identify at least one unsupported or overstated claim and revise the analysis.

**Evidence and check:** Collect the student's pre-AI reading, the verified matrix, the corrected interpretation, and a short explanation of which textual or precedential move changed the result. This assesses interpretation and source discipline, not the polish of the assistant's prose.

### A client-interview simulation with legally significant hidden facts

**Build:** Use a synthetic client file. Define the client's goals, emotional posture, chronology, facts the client volunteers, facts disclosed only after a well-formed question, and facts the client genuinely does not know. Add a faculty-only issue checklist. Instruct the assistant to stay in the role, answer one question at a time, avoid volunteering the whole story, and never score the student during the interview.

**Use in class:** In a Property class, the client might complain about a neighbor's security light; the hidden facts can map to the character of the harm, social utility, mitigation cost, hypersensitivity, and coming to the nuisance. In a clinic seminar, use a fictional intake involving missed notices, language access, or a deadline. Pairs conduct a ten-minute interview, then create a fact-to-element map before receiving the checklist.

**Evidence and check:** Debrief the sequence of questions, not merely whether every fact emerged. Require students to identify one leading question, one missed follow-up, and one fact that changed their legal theory. Never use a live client's facts unless the tool, account, consent, supervision, and data classification all permit it.

### A moot-court hot bench tied to the record

**Build:** Add a released moot-court record, the permitted authorities, competition rules, and the faculty rubric. Tell the assistant to act as one appellate judge at a time, ask a single concise question, ground follow-ups in the record or a cited authority, press on concessions, and avoid inventing facts. Create distinct judge profiles: jurisdiction, standard of review, remedy, and administrability.

**Use in class:** A student delivers an eight-minute argument while a partner records the questions. On the second round, switch the judge profile and side. Students tag each question as rule, record, standard of review, remedy, or policy and draft a tighter answer to the two questions they handled least effectively.

**Evidence and check:** Evaluate the student's oral responses and self-diagnosis against the human rubric. Faculty should spot-check every generated question against the closed record before using the assistant for a graded exercise.

### An AI-output audit for legal ethics and professional responsibility

**Build:** From a closed, public packet, create or save a polished but flawed AI memorandum. Seed realistic defects: a nonexistent citation, a case that does not support the proposition, omitted adverse authority, an overconfident conclusion, an undisclosed assumption, and an inappropriate request for confidential facts. Do not use an active client matter.

**Use in class:** Teams locate and classify each defect, verify every cited source, and prepare a corrected advice email plus a short supervision plan. Extend the discussion to competence, confidentiality, communication, candor, and reasonable review under ABA Formal Opinion 512.

**Evidence and check:** Grade the verification trail and professional judgment: what the student checked, where, what changed, what must be disclosed, and what still requires a lawyer. This directly rehearses the review work that persuasive AI prose can obscure.

## 2. Student use for learning

### A case-reading tutor that refuses to write the brief

**Build:** Give the assistant the assigned case and the professor's briefing headings. Require it to ask the student for procedural posture, legally material facts, issue, rule, application, and holding in that order. It should offer a small hint after an incomplete answer, quote or cite the assigned case when correcting the student, distinguish holding from dicta, and refuse requests for a finished brief until the student submits an attempt.

**Learning routine:** The student reads and briefs first, works through the tutor, then revises the brief in a different color or tracked version. For a case with a dissent, the tutor asks which premise the opinions actually dispute rather than merely summarizing both sides.

**Evidence and check:** Keep the original brief, two corrections with pinpoint support, and one unresolved question for class. The value is retrieval and self-correction, not a machine-produced case summary.

### Retrieval practice that adapts to a doctrine map

**Build:** Add the professor's learning objectives, an approved doctrine map, and short released examples. Instruct the assistant to ask one question at a time, require an answer and confidence rating before giving feedback, use hints before explanations, revisit missed concepts later, and keep a misconception log. Prohibit invented cases and new legal authorities.

**Learning routine:** A Civil Procedure student might cycle through subject-matter jurisdiction, personal jurisdiction, venue, Erie, and preclusion. The assistant varies the surface facts while keeping the tested concept stable. Once a week, the student writes a short explanation of the two most persistent errors without consulting the assistant.

**Evidence and check:** Use low-stakes pre- and post-checks written independently of the assistant. If the tutor marks a defensible legal answer wrong, the student flags it with the controlling source. That error report is part of AI literacy.

### A rule-synthesis and attack-outline coach

**Build:** Use the student's own notes plus assigned, verified cases. Tell the assistant not to produce a complete outline. Instead, it should ask the student to state a rule, identify the cases supporting each clause, reconcile tension, add exceptions, and test the rule against a new fact pattern.

**Learning routine:** The student submits a one-paragraph synthesis before asking for help. The coach returns a gap list such as “no authority for this limitation,” “two cases may use different levels of generality,” or “the remedy is missing.” The student checks the cases and revises the synthesis.

**Evidence and check:** Preserve three versions: the student's initial rule, the assistant's questions or gap list, and the source-verified final rule. Do not let the assistant become a substitute for reading the cases or building the outline.

### An exam issue-spotting coach built from released materials

**Build:** Add only professor-released practice questions, a simplified issue checklist or rubric, and model-answer excerpts authorized for study. Require the assistant to wait for the student's timed answer, ask the student to identify the legally significant facts, and give feedback by rubric category without rewriting the answer.

**Learning routine:** The student writes under a short time limit, self-scores, then uses the coach to locate missed issues, unsupported rule statements, absent counterarguments, and weak fact-to-rule links. A second answer must be written from a changed fact pattern without the assistant.

**Evidence and check:** Compare first and second performance on the same rubric. Do not upload a current take-home exam or use the coach where course policy prohibits AI assistance.

### A student-as-teacher legal-writing lab

**Build:** Ask the assistant to play a novice legal writer and produce a short response to a closed-universe problem. The assistant should make plausible organizational and analytical mistakes but use only the supplied sources. Give students the course rubric and require them to teach the assistant through specific feedback.

**Learning routine:** Students annotate the first draft, explain each revision request, and reject suggestions that would weaken the analysis. They then compare the revised AI draft with their own independent paragraph. This implements the “AI as student” approach: explaining and correcting another writer exposes the student's own understanding.

**Evidence and check:** Submit the initial output, annotations, feedback prompts, final output, and a reflection naming one place where the assistant resisted or misunderstood good feedback. Grade the student's diagnosis, not whether the AI eventually wrote an elegant paragraph.

## 3. Faculty use for productivity

### A course-maintenance impact map

**Build:** Create one private workspace containing the current syllabus, learning objectives, class calendar, assignment instructions, released rubrics, slide or lesson-plan text, and a document register with dates and controlling versions. Do not include identifiable student records. Instruct the assistant to distinguish quoted source content from suggestions and to cite the filename and section for every proposed change.

**Workflow:** When a case, rule, policy, or academic calendar changes, add the verified update and request an impact map: affected class session, learning objective, assigned reading, slide or handout, hypothetical, assessment, and student communication. The faculty member reviews originals and then updates each artifact.

**Quality control:** Use the map as a checklist, not an automatic amendment. A good assistant should say “no supplied source establishes this” when the update is incomplete.

### A class-preparation and hypothetical quality-control workspace

**Build:** Add the day's sources, the faculty member's existing class plan, learning objectives, and a short list of prior misconceptions with all student references removed. Ask for a timed run of show, Socratic branches, transition questions, and controlled factual variations that each change only one doctrinally material fact.

**Workflow:** Faculty select and edit the useful items, then ask the assistant to audit the final plan for unsupported propositions, accidental ambiguity, duplication, and whether every activity serves an objective. For a statutory course, request a separate check that quotations and section numbers match the supplied text.

**Quality control:** The professor remains the substantive editor. Retain the final human-approved plan as the next version's source rather than treating every generated alternative as institutional knowledge.

### An assessment and rubric alignment check

**Build:** Add the course learning objectives, draft assessment, grading rubric, timing assumptions, and one released sample response. Instruct the assistant to map each prompt to an observable skill, identify untested objectives, flag double-barreled or unintentionally ambiguous instructions, estimate the work required, and check whether the rubric rewards the intended analysis.

**Workflow:** For an exam, request an issue-coverage map and a fact-dependency check: which facts trigger which issues, which facts are noise, and whether two plausible readings create an unintended grading problem. For a seminar paper, compare milestone instructions, feedback opportunities, citation expectations, and final rubric criteria.

**Quality control:** Use the assistant to audit the instrument, not to make final grading decisions. Do not upload identifiable student submissions to an unapproved workflow, and never infer authorship or misconduct from AI-detection claims.

### A scholarship source-synthesis workspace

**Build:** Create one Project or assistant per article or bounded research question. Add the research plan, verified PDFs, the faculty member's notes, and a source register with full citations and status such as read, skimmed, or pending verification. Require page-level support, quotation marks for verbatim text, separate columns or sections for source claims and the assistant's synthesis, and an explicit “not found in supplied sources” response.

**Workflow:** Ask for a claim-evidence matrix, literature map, competing definitions, methodological differences, unresolved questions, and a counterargument list. After drafting, use the same workspace to trace each proposition back to the original and to build a response-to-comments matrix.

**Quality control:** Open and read the cited pages. Do not add confidential peer-review material, restricted datasets, or unpublished collaborator work unless the service and sharing settings are approved for that material.

### A committee or program decision log

**Build:** Use approved agendas, policies, public reports, and de-identified or appropriately classified meeting notes. Instruct the assistant to separate decisions, owners, deadlines, unresolved questions, evidence requested, and proposed language; preserve dates and document versions; and never fill a missing decision with an inference.

**Workflow:** Before a meeting, generate a source-linked briefing organized around decisions that must be made. After the meeting, a human owner confirms the action log. For a curriculum committee, the workspace might track how a proposed course maps to program outcomes, where coverage overlaps, what approvals remain, and which materials support each conclusion.

**Quality control:** Keep personnel evaluations, student petitions, accommodations, investigations, and privileged legal advice out of a general-purpose assistant unless an approved, access-controlled workflow has been established. Review access whenever committee membership changes.

### A controlled communications and course-FAQ assistant

**Build:** Ground the assistant only in the final syllabus, published course policies, assignment instructions, and approved administrative links. Instruct it to quote the controlling policy, distinguish a draft from a sent message, refuse to invent exceptions, and escalate individualized grading, accommodations, discipline, or legal questions to a person.

**Workflow:** Faculty can draft a weekly announcement, an assignment clarification, an office-hours recap, or a neutral FAQ update. The assistant can compare a proposed message with the controlling syllabus language and flag inconsistencies before the faculty member sends it.

**Quality control:** Keep the final human review visible. If students interact with the assistant directly, explain its role and limits, provide a non-AI route to the same information, and avoid collecting unnecessary student data.

## Research and governance basis

- Stanford Teaching Commons, Exploring the pedagogical uses of AI chatbots: presents seven roles adapted from Ethan and Lilach Mollick—mentor, tutor, coach, teammate, student, simulator, and tool—and cautions instructors and students to critically examine chatbot-generated feedback.
- Stanford Teaching Commons, Integrating AI into assignments: recommends beginning with meaningful assignments and observable learning objectives, aligning activities with evaluation criteria, assessing the student’s process, requiring robust citation when AI is used, and considering guided, low-stakes activities.
- John Bliss, *Teaching Law in the Age of Generative AI*, 64 Jurimetrics J. 111–61 (2024): provides law-school examples involving client simulations, factual variations, critique of AI-generated legal work, records of students’ AI interactions, tutoring, and flipped-classroom preparation.

- Greg Kestin et al., *AI Tutoring Outperforms In-Class Active Learning: An RCT Introducing a Novel Research-Based Design in an Authentic Educational Setting*, 15 Scientific Reports 17458 (2025): reports learning gains from a structured AI tutor that incorporated active engagement, sequential scaffolding, expert-crafted prompts and solutions, targeted feedback, and self-pacing. The authors caution that the findings may not extend to every discipline or to work requiring complex synthesis and higher-order critical thinking.
- American Bar Association,, Formal Opinion 512, *Generative Artificial Intelligence Tools* (2024): addresses lawyers' duties of competence, confidentiality, client communication, meritorious claims and contentions, candor toward tribunals, supervisory responsibilities, and reasonable fees when using generative AI.
- Stanford University IT, *Responsible AI at Stanford*: advises users not to enter Moderate or High Risk Data into a third-party AI platform or tool that is not covered by a Stanford Business Associates Agreement, to review Stanford-approved services by data-risk classification, and to recognize that AI outputs may be inaccurate or fabricated and cannot be trusted completely. If AI significantly influences a final product, Stanford advises users to consider disclosing how AI was used and citing it appropriately.

# Appendix A: A reusable build brief

Use this brief before creating any custom assistant or project.

## 1. Problem

- What recurring problem should this solve?
- Why is a saved setup better than a regular chat or a non-AI tool?

## 2. Audience

- Who will use it?
- What background knowledge can it assume?
- Who should not have access?

## 3. Inputs and sources

- What will users provide in each chat?
- What should persist as knowledge?
- Which sources control if documents conflict?
- Is every source approved for the tool, account, and audience?

## 4. Workflow

- What steps should the assistant follow?
- When should it ask a question?
- When should it stop or escalate to a person?

## 5. Output

- What sections, format, tone, and level of detail are useful?
- What evidence, citations, or uncertainty labels are required?

## 6. Boundaries

- What decisions require human judgment?
- What should never be treated as authoritative without verification?
- What sensitive information should users avoid providing?

## 7. Test plan

- What does success look like?
- Which ambiguous, unsupported, conflicting, adversarial, and privacy-sensitive cases will you test?

- Who will approve the setup before broader sharing?

## Appendix B: Source record

Research checked against official vendor and Stanford documentation on August 4, 2026.


OpenAI, Creating and editing GPTs
OpenAI, Sharing and publishing GPTs
OpenAI, File Uploads FAQ
OpenAI, Configuring actions in GPTs
Google, Use Gems in Gemini Apps
Google, Tips for creating custom Gems
Google, Share a Gem from Gemini Apps
Google, Upload and analyze files in Gemini Apps
Anthropic, How can I create and manage projects?
Anthropic, Upload files to Claude
Anthropic, Retrieval augmented generation for projects
Anthropic, Understanding Claude's personalization features
Stanford University IT, ChatGPT Edu at Stanford
Stanford University IT, Responsible AI at Stanford
Stanford Teaching Commons, Exploring the pedagogical uses of AI chatbots
Stanford Teaching Commons, Integrating AI into assignments
John Bliss, *Teaching Law in the Age of Generative AI*, 64 Jurimetrics 111-61 (2024)
Greg Kestin et al., *AI Tutoring Outperforms In-Class Active Learning: An RCT Introducing a Novel Research-Based Design in an Authentic Educational Setting*, 15 Scientific Reports 17458 (2025), https://doi.org/10.1038/s41598-025-97652-6
American Bar Association Standing Committee on Ethics and Professional Responsibility, Formal Opinion 512, *Generative Artificial Intelligence Tools* (July 29, 2024)