

ElevenLabs for Faculty

Content

[How to use this Information](#)

[Before you create an account](#)

[Rights, consent, disclosure, and accessibility](#)

[Set up ElevenCreative](#)

[Quick start: create a short reading roadmap](#)

[Quick start continued: choose a voice, model, and settings](#)

[Quick start continued: pronunciation, export, and publication](#)

[Choose and manage voices](#)

[Voice cloning: high caution](#)

[Build longer work in ElevenCreative Studio](#)

[Create a long-form audiobook](#)

[Dub faculty-created audio or video](#)

[Transcribe only authorized, Low Risk recordings](#)

[Generate sound effects and music sparingly](#)

[Generate images and video selectively](#)

[Real-world classroom uses for law faculty](#)

[More law-school teaching designs](#)

[Design a student assignment that preserves learning](#)

[Faculty learning and productivity uses](#)

[Quality-control checklist](#)

[Troubleshooting](#)

[Official sources and verification notes](#)

[Appendix A: Sample Text-to-Speech Script](#)

[Appendix B: Video Recordings](#)

STANFORD USE STATUS — NOT APPROVED ElevenLabs is not currently a Stanford-approved AI service. Use only public or appropriately authorized Low Risk material. Do not upload student records, client or privileged information, research data, personnel information, protected health information, nonpublic financial information, credentials, unpublished confidential scholarship, or any voice or recording you lack permission to use.

Research verified August 4, 2026

How to use this Information

This tutorial is written for faculty and instructional teams who want to experiment with AI-generated speech and audio without treating the tool as an authority, an approved data environment, or a substitute for human teaching judgment. The primary example is a faculty-authored, ninety-second audio reading roadmap. Later sections cover longer projects, dubbing, transcription, sound effects, music, and law-school teaching designs.

By the end, you should be able to

- decide whether an audio task is appropriate for an unapproved third-party service;
- generate, review, and download a short text-to-speech recording;
- choose between a library voice, a designed synthetic voice, and a voice clone;
- build and export a longer audio or video project in ElevenCreative Studio;
- use dubbing, transcription, sound effects, and music with appropriate human review;
- design law-school activities in which the audio supports an observable learning objective; and
- publish accessible audio with a transcript or accurate captions and a clear AI disclosure.

What ElevenLabs is

ElevenLabs describes its platform as AI voice infrastructure for text-to-speech, speech-to-text, voice cloning, conversational agents, and generative audio. Its no-code browser environment is now called ElevenCreative. This guide concentrates on the creator-facing workflows that a faculty member can operate in a web browser.

Capability	What it does	Useful faculty example	Primary caution
Text to Speech	Turns a prepared script into synthetic narration.	A short reading roadmap or course orientation.	Pronunciation and emphasis still require review.
Studio	Combines narration, captions, music, sound effects, and video on a timeline.	A multi-part evidence simulation with a transcript.	Imports and AI-generated scripts can alter source material.

Capability	What it does	Useful faculty example	Primary caution
Dubbing	Translates and revoices audio or video.	A multilingual version of a faculty-created public explainer.	Legal terms and translation need fluent human review.
Speech to Text	Transcribes audio or video and separates speakers.	A transcript of a public lecture you own.	Do not upload class, client, or meeting recordings.
Voices	Provides library, designed, instant-cloned, and professionally cloned voices.	A synthetic fictional client voice.	Never imitate a real person without rights and consent.
Sound Effects / Music	Generates audio from prompts.	A neutral scene cue or short intro bed.	Licensing, disclosure, and evidentiary meaning matter.
Audiobooks	Turns a manuscript into chaptered long-form narration, with export and publishing workflows.	An audio edition of faculty-authored course notes, an open educational resource, or published scholarship.	Book-length generation can consume substantial credits; rights, transcript access, and distribution terms still require review.
Image & Video	Generates images and short video with multiple visual models; can also apply lip sync and other media enhancements.	A disclosed talking visual for a faculty-authored explainer or a synthetic classroom character.	Visual generation uses the same credit pool. Prefer Stanford-provided tools for ordinary image/video generation; use ElevenLabs when audio integration or lip sync is the reason.

ElevenAgents, API integrations, telephony, and live conversational bots are outside this guide. Consult Stanford's Information Security Office before any Stanford deployment.

Before you create an account

1. Treat Stanford approval as a boundary, not a disclaimer

STOP BEFORE UPLOAD ElevenLabs is not on Stanford's current Approved AI Services list. An opt-out setting, a paid plan, or a vendor privacy statement does not make the service Stanford-approved.

Stanford's Responsible AI guidance explains that data entered into third-party AI services is transmitted to and stored on external servers. For tools not covered by a Stanford agreement, avoid Moderate or High Risk Data. Even with Low Risk Data, ask whether you would be comfortable with the material becoming public.

Never upload

- **Student information:** names tied to work, grades, accommodations, advising notes, education records, or identifiable class recordings.
- **Legal practice material:** real client facts, privileged communications, confidential files, work product, interview recordings, or litigation strategy.


- **Research and employment material:** human-subject data, unpublished confidential drafts, peer-review material, personnel information, or information protected by an NDA.
- **Security and identity material:** credentials, access tokens, financial information, health information, or biometric/voice data from another person.
- **Third-party content:** casebook chapters, paid databases, recordings, music, or other works unless your use is authorized.

Stanford guidance: [Responsible AI at Stanford](#)

Current service matrix: [Stanford Approved AI Services](#)

2. Review ElevenLabs' data-use setting

ElevenLabs states that it uses certain submitted data to improve its audio models unless the user opts out. To change the setting, open your profile, choose Terms and privacy, open Data use, disable Improve the models for everyone, and save the change. ElevenLabs says this stops new data from being used to train its models. It does not remove other processing needed to provide, secure, or moderate the service, and it does not change Stanford's approval status.



Turn off model-improvement data sharing before experimenting with public or Low Risk content.

READ: [Is my data used to improve ElevenLabs' AI models](#)

3. Understand the account and plan terms

- This is not a Stanford-managed account. Do not assume Stanford support, procurement review, SSO, retention controls, or institutional contracting.
- ElevenLabs' Terms of Service, last updated March 31, 2026, require users engaging in noncommercial use to use a personal email address. The Terms also say that ElevenLabs may treat the organization associated with a business email address as the controller of personal data in the account, and that users acting on behalf of an organization must have authority to bind it. Recheck the Terms before registering, using a Stanford email address, or representing that an account is affiliated with Stanford.

- Features, credits, quality, export formats, attribution requirements, and commercial rights vary by plan and may change.
- The free plan currently requires attribution for shared noncommercial output and does not include a commercial license. Paid-plan rights remain subject to the Terms, use policy, beta terms, and your rights in all inputs.
- We recommend that you do not require students to purchase a plan or open an account without an equivalent, no-cost, non-AI way to meet the same learning objective.

See More Details:

- Vendor terms:
ElevenLabs. (2026, March 31). *Terms of service*. Retrieved August 5, 2026, from <https://elevenlabs.io/terms-of-use>
- Plan rights and attribution: [Publishing ElevenLabs-generated content](#)
- Current plan details: [ElevenLabs pricing](#)

Rights, consent, disclosure, and accessibility

Use a voice you are entitled to use

For teaching, the safest default is a stock/library voice or a wholly synthetic voice created with Voice Design. If you clone a voice, clone only your own. Do not clone or imitate a student, client, witness, judge, colleague, public figure, actor, or deceased person. Do not use a synthetic voice to create the impression that a real person said something they did not say.

VOICE-CLONING RULE Professional Voice Cloning can be created only for your own voice and requires verification. If a collaborator wants to share a Professional Voice Clone, they must create and verify it in their own account and share it with you. For classroom simulations, prefer a designed synthetic voice rather than a clone.

ElevenLabs' Prohibited Use Policy bars unauthorized, deceptive, or harmful impersonation and bars using a replicated voice without consent or legal right. It also requires a qualified professional in the loop and clear disclosure for tailored legal services. Classroom content must not be presented as individualized legal advice.

- Vendor policy: [ElevenLabs Prohibited Use Policy](#)
- Vendor rule for Professional Voice Clones: [Whose voice can be cloned](#)

Disclose synthetic audio

Add a disclosure when audio is AI-generated, translated, altered, or performed by a synthetic voice—especially when the result could be mistaken for a real person or authentic evidence. Put the disclosure in the audio, beside the player, and in the transcript when practical.

Suggested course-media disclosure:

This audio uses an AI-generated voice. The script was written and reviewed by [faculty/teaching team]. It is a teaching aid, not an authoritative statement of law or legal advice.

Suggested simulation disclosure:

This is a fictional classroom simulation. The people, facts, and voices are synthetic and do not represent a real client, witness, judge, or lawyer.

Publish the text equivalent at the same time

Stanford's Office of Digital Accessibility states that audio-only content needs an accurate transcript and video needs accurate synchronized captions. Auto-generated captions and transcripts must be reviewed for punctuation, names, technical terms, speaker identification, and meaningful sound cues. Audio should supplement—not replace—the accessible text, readings, citations, or instructions students need.

- For audio-only files, provide an accurate transcript adjacent to the player or download link.
- For video, provide accurate synchronized captions; a transcript alone is not a substitute for captions.
- Include relevant non-speech cues such as [door closes], [phone vibrates], or [music fades].
- We recommend not grading students on production polish, accent, vocal affect, or use of a paid synthetic voice unless those qualities are themselves an explicit and equitable learning objective.

Stanford accessibility guidance: [Audio and Video](#)


Stanford captioning details: [Captioning Basics](#)

Preflight checklist

Item	Proceed only if the answer is yes
Data	The script and any upload contain only public or properly authorized Low Risk content.
Rights	You may use every text, recording, image, voice, music reference, and sound included.
Consent	Any identifiable speaker agreed to the specific recording, transformation, and sharing.
Accuracy	A qualified person will review legal claims, names, quotations, translations, and citations.
Disclosure	The audience will know when voice, translation, or other material is AI-generated.
Access	An accurate transcript or captions and a non-AI pathway are available.

Set up ElevenCreative

1. Open the ElevenLabs website and review the current Terms, Privacy Policy, Prohibited Use Policy, and pricing before creating an account.
2. Use a personal account for personal, noncommercial experimentation if that is what the current Terms require. Do not treat the account as Stanford-managed.
3. Open profile settings and turn off Improve the models for everyone before entering any content.
4. From the left navigation, identify tools such as Text to Speech, Speech to Text, Voice Changer, Sound Effects, Image & Video, Music, Audiobooks, Studio, Dubbing, and Voices.
5. Confirm your current plan, credits, export options, and attribution obligations before committing to an assignment or publication schedule.



ElevenCreative groups browser-based speech, voice, Studio, dubbing, transcription, music, and sound tools in one workspace.

Choose the smallest tool that fits

If you need to...	Start with...	Why
Narrate one short script	Text to Speech	Fastest path; easiest to review line by line.
Build multi-part audio or combine media	Studio	Timeline editing, multiple tracks, captions, and exports.
Translate existing audio/video	Dubbing	Speaker detection and translated synthetic speech.

If you need to...	Start with...	Why
Create a transcript from an authorized recording	Speech to Text	Language detection, timestamps, speakers, and export. Note: Google Gemini and Gemini Notebook can also handle this task well in a Stanford approved environment.
Invent a fictional voice	Voice Design	Avoids imitating a real person.
Create a short nonverbal cue	Sound Effects	Text-prompted variations and loop controls.
Turn a long faculty-authored text into chaptered audio	Audiobooks	Purpose-built workflow for long-form narration, chapter organization, export, and optional ElevenReader distribution.
Create a standalone image or general-purpose AI video	Stanford Gemini / ChatGPT Edu first	Stanford-provided services cover common image generation, and Gemini Enterprise includes limited video generation. Reserve ElevenLabs visual credits for workflows that benefit from its audio and lip-sync integration.

GOOD DEFAULT For a first faculty project, use Text to Speech with a library or designed synthetic voice. Avoid voice cloning, dubbing a real person, or uploading recordings until you can document rights, consent, data classification, and review.

Quick start: create a short reading roadmap

This workflow produces a ninety-second audio orientation that prepares students to read a case without giving them the answer. The script is faculty-authored, the voice is synthetic, and the final audio is published with the reviewed script as its transcript.

Step 1: Write for listening

- State the instructional purpose in the opening sentence.
- Use short sentences and one idea per sentence. Write numbers, abbreviations, and symbols the way they should be spoken.
- Name the authority and assignment precisely, but do not paste copyrighted casebook commentary into the tool.
- Include two or three listening questions. Do not summarize away the reasoning students are expected to do.
- End with the AI disclosure and a direction to consult the assigned sources.


Sample script — replace bracketed text before use

This short audio is a reading roadmap for [case name and citation]. As you read, identify the procedural posture, the legal question the court actually had to decide, and the facts the court treated as legally significant. Mark the sentence or paragraph that best states the court's rule. Then ask whether each disputed statement is necessary to the judgment or is better understood as dicta. Bring one passage that supports your answer and one uncertainty to class. This audio uses an AI-generated voice. The script was written and reviewed by [name]. Consult the assigned case and course materials for the controlling text.

LEGAL ACCURACY Text-to-speech creates a performance of your script; it does not verify the script. Confirm every name, citation, quotation, doctrinal statement, and pronunciation against the source before and after generation.

Step 2: Open Text to Speech and enter the script

1. Select Text to Speech in the left navigation.
2. Paste the reviewed script into the text box. Remove comments, footnote markers, raw URLs, unusual symbols, and text that should not be spoken.
3. If the interface displays a character limit, divide longer material at a logical paragraph or section boundary rather than mid-sentence.




Paste a reviewed script into Text to Speech, then choose the voice and model before generating.

Quick start continued: choose a voice, model, and settings


Step 3: Choose a voice

Voice choice has more effect on the result than the fine-grained settings. For a reading roadmap, begin with a calm, clear library voice or a designed synthetic voice. Test the voice with names and legal terms from your script. A native voice for the target language usually produces a more natural accent.



Step 4: Choose a model

Use a standard/high-quality model for recorded course media and a multilingual model when the script is not exclusively English. ElevenLabs describes Flash models as optimized for low latency and Multilingual v2 as stable for long-form generation. Model names and language coverage change, so choose based on the current interface and test the selected voice-model pairing.



Step 5: Adjust only what you can hear

The screenshot shows a settings panel with the following controls:

- Speed:** A slider ranging from "Slower" to "Faster". The marker is positioned approximately two-thirds of the way towards "Faster".
- Stability:** A slider ranging from "More variable" to "More stable". The marker is positioned approximately two-thirds of the way towards "More stable".
- Similarity:** A slider ranging from "Low" to "High". The marker is positioned approximately two-thirds of the way towards "High".
- Language Override:** A toggle switch currently turned off.
- Output Format:** A dropdown menu showing "MP3 44.1 kHz (128kbps)".

Voice and model choice matter more than small slider changes; adjust settings only after comparing audible results.

Control	What it changes	Practical starting point
Speed	Slows or accelerates speech when the selected model supports it.	Stay near the default; adjust in small steps after listening.
Stability	Balances consistency against expressive variation.	Moderate to high for instructional narration; lower only if the result sounds flat.
Similarity	How closely a voice model follows its source characteristics.	Leave near default; high values can reproduce source artifacts.
Style exaggeration	Amplifies the source voice's style but can reduce stability.	Keep at zero unless a tested pedagogical reason requires it.
Speaker Boost	Subtly increases similarity on supported models.	Optional; compare before spending extra processing time.

Vendor reference: [Text to Speech product guide](#)

Step 6: Generate, listen, and compare

1. Select Generate Speech and listen through headphones and ordinary speakers.
2. Check names, case citations, section numbers, Latin phrases, acronyms, dates, and homographs in context.
3. Change one variable at a time—voice, model, punctuation, speed, or stability—so you can identify what improved the result.
4. Regenerate only the line or section that needs repair when the interface permits. Keep a final script that matches the exported audio.

Our Sample

1. View the sample text (see Appendix A)
2. [Watch a screen-recording of text-to-speech in action.](#)
3. [Listen to the Generative Version A](#) (Liam Voice)
4. [Listen to Generative Version B](#) (Marsh voice)

Quick start continued: pronunciation, export, and publication

Fix pronunciation before adding performance effects

Problem	Try this first	Example
Digits or symbols sound wrong	Write them as spoken words.	“section two thirty” instead of “§ 230.”
Abbreviation is misread	Spell it out or add periods/spaces.	“Federal Rules of Evidence” before using “F R E.”
Name is mispronounced	Use a phonetic respelling in the spoken copy or a pronunciation dictionary if available.	Keep the published transcript in standard spelling.
Pacing is too dense	Shorten the sentence and add punctuation or a supported break.	Use pauses sparingly; excessive break markup can create artifacts.
Voice becomes unstable	Reduce style exaggeration, restore default settings, or use a more stable voice/model.	Generate a short test before the full script.

Export and preserve the production record

- Download the final file in the format appropriate for your course platform. MP3 is usually the simplest for distribution; WAV is useful for further editing.
- Save the exact final script, source list, voice name or identifier, model, settings, date, plan/attribution status, and disclosure text.
- Rename the files descriptively, for example: Evidence_Week03_Reading_Roadmap_v1.mp3 and Evidence_Week03_Reading_Roadmap_transcript.docx.
- Do not rely on platform history as the only copy. ElevenLabs does not guarantee perpetual availability of content under all plans.

Publish accessibly

- Place the disclosure next to the audio player and include it in the transcript.
- Upload the reviewed transcript at the same time as the audio. If the audio is inside video, add accurate synchronized captions.
- Keep the original reading, authorities, and citations visible. The audio is an orientation or alternative modality, not the source of law.


- Invite students to report pronunciation, caption, or access problems and correct the source file rather than explaining errors only in class.

Choose and manage voices

Voice type	Best fit	Quality / control	Faculty guidance
Default or Voice Library	Fast narration with an existing voice.	Quality varies; filters and previews help.	Preferred starting point. Confirm current sharing and usage terms.
Voice Design	A fictional voice described in words.	Fast exploration; results vary with prompt and preview text.	Preferred for fictional clients, witnesses, or narrators.
Instant Voice Clone	A quick representation from short samples.	Less consistent than a professional clone.	Use only for your own voice; avoid for classroom characters.
Professional Voice Clone	A high-fidelity model of your own voice.	More training audio and verification; plan restrictions apply.	Only your own voice. A collaborator must verify and share theirs.

Find a Voice Library voice

1. Open Voices and choose the Voice Library.
2. Filter by language, accent, use category, quality, and other available attributes.
3. Preview several voices with the same short portion of your actual script.
4. Choose a voice for clarity and consistency, not for stereotyped associations with identity, authority, nationality, gender, age, disability, or social status.
5. Save or select the voice, then test it again in Text to Speech before producing the full audio.



Compare voices using the same script and choose for clarity, stability, and fit—not identity stereotypes.


Vendor reference: [Voice Library guide](#)


Create a fictional voice with Voice Design

Use Voice Design when no library voice fits and you need a voice that is explicitly fictional. Describe audible qualities relevant to the performance—language, accent, pace, tone, age range, vocal texture, and delivery—without naming or asking to imitate a real person.

Example voice-design prompt

A clear, measured adult voice for a fictional client voicemail in an educational simulation. Contemporary American English, moderate pace, thoughtful but slightly uncertain delivery, natural pauses, no resemblance to a public figure or named person.





Voice Design creates synthetic voices from descriptive prompts; do not request an imitation of a real person.

Vendor reference: [Voice Design guide](#)

Voice cloning: high caution

PREFERRED PRACTICE Use a library or designed synthetic voice for course simulations. Clone only your own voice when a clear teaching or accessibility need outweighs the risks and you understand the current terms.

If you clone your own voice

- Record only your voice in a quiet room. Remove other speakers, background media, and confidential conversation.
- Read the current voice-processing notice and Terms before uploading. Voice data may be treated as biometric data in some jurisdictions.
- Start with a short, disposable public script and test privately. Do not publish the clone to the Voice Library unless you intentionally want public discovery and have reviewed the sharing terms.
- Use the account's verification process for a Professional Voice Clone. Do not evade or work around voice verification.
- Label generated audio as synthetic and preserve a way to withdraw or replace course media if the voice model is compromised or terms change.
- Delete unused clones and recordings from the platform where available, recognizing that deletion and training opt-out have limits described in the vendor's policies.

Instant Voice Clone

Clone your voice with only 10 seconds of audio.

2 minutes

Professional Voice Clone

Create the most realistic digital replica of your voice. Requires at least 30 minutes of clean audio.

5 minutes

You need to be on at least the Creator plan to use Professional Voice Cloning

Subscribe

Avoid noisy environments

Background sounds interfere with recording quality results.

Check microphone quality

Try external units or headphone mics for better audio capture.

Use consistent equipment

Don't change recording equipment between samples.

Click to upload, or drag and drop
Audio or video files up to 10MB each

or

Record audio

☐ 10 seconds of audio required

Next

Use a designed voice for fictional characters; Professional Voice Cloning is for the verified account holder's own voice.

Do not use cloning for

- a simulated statement by a real judge, lawyer, client, witness, politician, celebrity, colleague, or student;
- a memorial or historical reenactment that imitates a deceased person without documented rights and careful disclosure;
- a prank, assessment trap, undisclosed authenticity exercise, or communication that could leave the classroom and mislead others;
- automated legal guidance, intake, advising, grading, discipline, accommodation, or other high-stakes decision; or

- any recording obtained without the speaker’s informed agreement to the recording and the proposed AI transformation.


Vendor privacy notice: [ElevenLabs Privacy Policy](#)

Vendor explanation: [Voice cloning: how it works](#)


Build longer work in ElevenCreative Studio

Studio is the current end-to-end timeline editor. It can combine narration, captions, video, music, and sound effects; support sentence-level timing; collect comments; and export audio or video. Older tutorials may refer to “Projects” or “Voiceover Studio.” Use the current Studio interface instead.

Start the project




1. Open Studio and choose Upload or New blank project. We recommend using the video project exclusively as it works well with audio and has more features including an editing timeline.



2. Upload only content you are permitted to process. Studio can analyze text, audio, or video and choose a layout. Inspect every imported heading, paragraph, caption, and media track.
3. For maximum control, start with a blank audio project and paste the reviewed script in sections. Add only the tracks the learning objective requires.

- Choose the model and project settings. Generate a short test section before committing credits to the full project.



Start from a reviewed script or authorized file; inspect any automated import or generated script before narration.

Watch our brief, ElevenLabs, studio demo: [View Video](#) (4 minutes, 37 seconds)

Faculty scenario to try creating: a three-part evidence simulation

Track / section	Faculty builds	Students do
1. Disclosed client voicemail	A fictional intake message performed by a designed synthetic voice.	List facts, ambiguities, and follow-up interview questions.
2. Environmental sound cue	A clearly labeled synthetic door-close or phone vibration, if relevant.	Explain whether the cue changes an inference and what foundation would be needed.
3. Faculty debrief	A separate narration identifying the learning objective and source authorities.	Compare their initial inference with the debrief; preserve disagreement.
Transcript / captions	Accurate text, speaker labels, disclosures, and sound cues.	Use the transcript to cite exact language in analysis.

DESIGN RULE Do not let production realism outrun instructional transparency. Students should know the exercise is synthetic before using it unless an IRB-approved research design or other formal process explicitly governs the activity. A classroom “gotcha” can undermine trust and may circulate outside its context.

Edit, review, and export a Studio project

Work section by section

- Split the script at logical points—introduction, source audio, questions, debrief—not at arbitrary character counts.
- Assign one voice per fictional speaker and maintain a speaker log. Avoid switching identity-coded voices merely to signal credibility or status.
- Regenerate the smallest faulty sentence. Recheck timing after any text edit because fixed durations can make speech unnaturally fast or slow.
- Keep music and sound effects below speech. Remove them if they interfere with comprehension or captions.
- Review captions against the final audio, not an earlier script. Correct names, citations, legal terminology, and non-speech cues.

Be cautious with AI-generated scripts

Studio can generate a podcast or script from a document, URL, or prompt. That workflow creates a new script; it is not a faithful audio rendering of the source. For legal teaching, compare every claim with the cited authority, restore omitted qualifications, remove invented quotations or citations, and label the result as a faculty-reviewed teaching draft. If you need verbatim narration, start from your own finalized script rather than a generative podcast workflow.

Export package


- final audio or video file;
- accurate transcript or synchronized captions;
- source and rights log;
- AI disclosure and voice/translation disclosure;
- production record containing voice, model, settings, and review date; and
- a text-only equivalent assignment or activity path.

Currently, the Eleven Lab exports will have a watermark and ElevenLabs identifier add the end of each video you generate.

Vendor reference: [ElevenCreative Studio guide](#)

Create a long-form audiobook


Audiobooks is a dedicated long-form workflow for turning an authorized manuscript into organized narration. It is likely to be especially useful to faculty who already have a polished text and want a listening edition without building every chapter manually in Studio. ElevenLabs supports fixed, pre-generated narration as well as dynamic narration for ElevenReader, where listeners can choose a preferred narrator voice.



The Audiobooks bookshelf separates audiobook creation from the general Studio workflow and also offers direct publishing to ElevenReader.

Why faculty may care


- Create an audio edition of faculty-authored course notes, a public or openly licensed teaching resource, or already-published scholarship.
- Organize a longer work by chapter and keep narration, pronunciation, and export in one project rather than assembling many short Text to Speech files.
- Use a single narrator for a conventional reading, or multi-cast narration when distinct fictional characters genuinely support the teaching design.
- Publish a reviewed audiobook to ElevenReader or export audio for another authorized distribution channel. Dynamic narration can give listeners control over narrator choice.



The onboarding screen distinguishes crafted audiobooks from free dynamic narration and highlights publishing and distribution.

Basic workflow

1. Open Audiobooks from the left navigation and choose Create an Audiobook.
2. Upload an authorized EPUB or PDF in the current creation flow. Do not upload a casebook, licensed database export, student work, client material, or confidential draft unless you have a documented right and an approved data environment for that content.
3. Choose Single cast for a conventional audiobook. Use Multi cast only when multiple synthetic voices add a clear instructional benefit, and review every character-to-voice assignment.
4. Choose the narration model, inspect the imported structure, and generate a short sample before committing to a full chapter. Fix formatting and pronunciation while the test is still small.
5. Review chapter boundaries, names, citations, quotations, and any language that could change meaning when heard rather than read. Keep the final text available as the transcript or equivalent accessible version.
6. Export the finished book or publish only after checking current rights, attribution, payout, and distribution terms. If you use dynamic narration, remember that music, sound effects, and external audio are not part of dynamic playback.



The current setup asks for a manuscript, narration style, and model before generation. Multi-cast narration is available alongside the simpler single-cast option.


COST CONTROL Audiobook generation uses the shared ElevenLabs credit pool. Generate one representative chapter or several paragraphs first, then estimate the full project before committing. Playback of audio that has already been generated does not spend additional generation credits, but generating uncreated sections during playback or export does. Long-form projects can consume a large share of a monthly plan.

ACCESSIBILITY Treat the audiobook as an additional modality, not the only copy. Publish the reviewed text at the same time, preserve headings and citations in the text version, and disclose synthetic narration.

Vendor reference: [ElevenLabs Audiobooks documentation](#)


Dub faculty-created audio or video


Dubbing translates speech and generates new audio while attempting to preserve speaker characteristics, timing, tone, and emotion. That can make a translation sound authoritative even when a legal term, name, quotation, or cultural reference is wrong. Use dubbing for material you created or are authorized to transform, and require review by someone fluent in the target language and the legal subject.



Basic workflow

1. Open Dubbing from the navigation or start a dub from Studio.
2. Upload an authorized audio/video file or use an authorized URL. Do not upload class discussion, client interviews, office hours, or copyrighted recordings you lack permission to transform.
3. Select the source and target languages.
4. In the advanced section, you may control how close or different the dubbed voice should sound to the original voice.







5. Generate a short test. Compare the transcript and translation line by line before completing the full project.

6. Correct terminology, pacing, speaker assignment, captions, and pronunciation. Export the source-language transcript, reviewed translation, dubbed media, and captions together.

A NOTE ABOUT COSTS: Dubbing is an expensive feature. For a video less than 5 minutes and dubbing in two languages, the credit expenditure is 124,813 credits and if you don't like the results, there is no way to get a refund.



Compare this one time generation costs to these plans below as of August 6, 2026):



Legal-language review protocol

Check your dubs for these:

Review layer	Reviewer checks
Text fidelity	No facts, hedges, dates, quotations, citations, or speaker turns were added, removed, or reassigned.
Legal terminology	Terms of art, institutional names, case names, statutes, and procedural labels use an accepted target-language formulation.
Audience	The translation does not become individualized legal advice and clearly identifies the governing jurisdiction and limitations.
Audio performance	Pacing, emphasis, emotion, and pronunciation do not change the apparent certainty or meaning.
Access and disclosure	Reviewed captions/transcript accompany the dub, and the synthetic translation is disclosed.

DO NOT ASSUME A fluent-sounding dub is not proof of an accurate translation. For legal information, a bilingual subject-matter reviewer—not the model—must approve the final wording.

Vendor reference: [Dubbing overview](#)


Transcribe only authorized, Low Risk recordings

NOT FOR CLASS OR CLIENT RECORDINGS Because ElevenLabs is not Stanford-approved, do not use its Speech to Text tool for classroom discussion, office hours, meetings, interviews, advising, clinics, client work, or human-subject research. A participant’s permission to be recorded does not by itself approve processing in this service.


Appropriate experiments are limited to recordings that are public, owned by you, or otherwise authorized and contain no sensitive Stanford data—for example, your own published public lecture or a fully fictional simulation.

Basic workflow

1. Open Speech to Text and choose Transcribe files.



2. Upload the authorized audio or video. Select the known language or leave detection on when appropriate. In addition to uploading, you may also choose to record yourself, list a YouTube video, or another media URL.



3. Enable audio-event tags if laughter, applause, music, or other sounds matter. Add key terms such as case names, statutes, Latin phrases, and speaker names to improve recognition where the current interface permits.
4. Open the result, rename speakers, play audio from suspect words, and edit the transcript.
5. Export the needed format and conduct a final comparison against the audio before using the transcript as captions, quotations, evidence, or research data.

Human review is mandatory

- A word-level timestamp is not proof that the word is correct.
- Speaker diarization can assign words to the wrong speaker, especially with overlap or poor audio.
- Automatic captions require editing before publication; verify capitalization, punctuation, names, citations, and meaningful sound cues.
- Do not quote a transcript in scholarship, a filing, or teaching material without returning to the source recording and any applicable citation rules.


Vendor reference: [Speech to Text product guide](#)

Generate sound effects and music sparingly

Sound Effects

Sound Effects creates several short audio variations from a text description and can control duration, looping, and prompt influence. In law-school media, a sound cue should serve an explicit learning purpose—not merely make a simulation more dramatic.

1. Open Sound Effects and describe one sound at a time, such as “single office door closes softly, clean recording, no voices.”
2. Choose a short duration and generate variations. Listen for unwanted speech, music, or events.
3. If the sound is part of an evidence exercise, identify it as synthetic and explain whether students should treat it as a stipulated fact, a disputed inference, or irrelevant atmosphere.
4. Download the selected file and include the sound cue in the transcript, for example [office door closes].



Generate a short, purpose-built sound cue, review all variations, and disclose the cue as synthetic when it affects interpretation.


Vendor reference: [Sound Effects guide](#)

Music

Eleven Music can generate and edit songs from prompts and, in some plans, use a short audio reference. Avoid style prompts that identify living artists, copyrighted songs, or recordings you do not control. Check the current Music Terms and plan-specific rights before public distribution.

Note: other platforms to consider for music generation include [Suno](#) and [Gemini's Lyria3 model](#). Suno is currently the lead in music generation innovation, but Gemini's music generator can be used within the gemini app.

- Use instrumental music only when it supports orientation, transitions, or public outreach; keep it below narration.
- Do not use music to manipulate perceived credibility, emotion, or seriousness in a client, witness, or evidence simulation unless that effect is the stated object of analysis.
- Keep a clean version without music when it improves comprehension or accessibility.
- Document the prompt, plan, generation date, and applicable license/attribution terms.



Use only authorized prompts and references, then verify the current plan's licensing and attribution terms before publication.

Vendor reference: [Eleven Music guide](#)

Generate images and video selectively

ElevenCreative now includes an Image & Video workspace with access to multiple image and video models. It can generate new visuals, use image or video references, upscale media, and route finished assets into Studio. The distinctive reason to use it in a faculty workflow is not basic image generation; it is the ability to connect visuals with ElevenLabs audio, especially lip sync and talking-character workflows.




Image & Video supports image, video, and lip-sync workflows and displays the model and generation settings before submission.

Recommended platform choice

PREFERRED PRACTICE For ordinary image generation, start with Stanford-provided [Google Gemini](#) or [ChatGPT Edu](#) rather than spending ElevenLabs credits.

For general [AI video](#), start with Stanford Gemini Enterprise / [Google Vids](#)* when the available daily quota is sufficient. Use ElevenLabs Image & Video when the project specifically benefits from its audio stack, such as synchronizing a synthetic voice to a permitted image or video.

Images: Gemini and ChatGPT Edu are strong default choices for concept art, diagrams, illustrative course graphics, and iterative image editing in Stanford-provided environments.

Video: Gemini Enterprise includes limited video generation. Stanford currently lists one video request per day for Standard and three per day for Pro, subject to change; check the current service page before planning an assignment.

Lip sync: ElevenLabs is more distinctive when you already have authorized audio and an authorized image/video and need the mouth movement synchronized to the new speech. This can be useful for a disclosed fictional character or a faculty-owned explainer, but it should not be used to make a real person appear to say something they did not say.

Cost: ElevenLabs Image & Video uses the same shared credits as speech, dubbing, music, and audiobooks. Video generation is a paid-plan feature, and the cost varies with model, duration,

resolution, and number of variations. The interface shows the credit cost before generation; treat that preview as a budget checkpoint.

Iteration: Re-running or recreating a generation consumes credits again. Draft the prompt and decide the aspect ratio, duration, and reference media before submitting expensive video or lip-sync jobs.

*AI features in Google Vids are currently disabled by UIT in Stanford Gemini Enterprise

Faculty use cases

- Create a neutral, clearly synthetic talking character to introduce a fictional fact pattern.
- Turn a faculty-authored audio abstract into a short captioned visual explainer using visuals you own or are permitted to transform.
- Create illustrative still images for a lecture or simulation, then move the selected asset into Studio for narration and captions.
- Demonstrate synthetic-media literacy by comparing an original clip with a disclosed lip-synced version and asking students what authentication evidence would be needed.

RIGHTS AND DISCLOSURE Use only images, video, audio, and likenesses you are entitled to transform. Do not upload student, client, witness, colleague, or other identifiable-person media to create a talking avatar or altered performance. Label synthetic or materially altered visuals, and provide accurate synchronized captions for any video with speech.

Vendor reference: [ElevenLabs Image & Video documentation](#)

Real-world classroom uses for law faculty

The strongest uses begin with an observable legal-learning objective and make the student—not the synthetic audio—perform the legal reasoning. Each design below can be built with fictional facts and public or faculty-authored material.

1. Fictional client intake: facts, uncertainty, and interviewing

Component	Design
Objective	Distinguish facts from assumptions and formulate sequenced, non-leading follow-up questions.
Build	Create a ninety-second voicemail using a designed fictional voice. Include four legally relevant facts, two ambiguities, one emotionally salient but legally peripheral detail, and one timeline gap.
Student work	Submit a fact/assumption table, eight interview questions, and a short explanation of question order.
Debrief	Reveal the faculty-authored fact map and compare how vocal tone changed perceived credibility or urgency.

Component	Design
Guardrails	No real client facts or voice. Disclose the simulation. Provide transcript and a text-only option. Do not use AI to score students.

2. Evidence and deepfake authentication lab

Component	Design
Objective	Develop a Rule 901 authentication theory and identify what additional evidence is needed before admitting or relying on an audio file.
Build	Provide a disclosed synthetic recording, its production log, and a deliberately incomplete chain-of-custody packet. Do not imitate a real person.
Student work	Draft foundation questions, identify alternative explanations, and separate authenticity, hearsay, best-evidence, and prejudice issues.
Debrief	Compare arguments under traditional evidence rules with emerging deepfake risks. Inspect metadata and the synthetic-audio disclosure.
Guardrails	The recording stays inside the course site; filename and opening audio label it synthetic; no political candidate or public-figure voice.

3. Oral-argument bench question bank

Component	Design
Objective	Answer difficult questions directly, preserve a theory of the case, and use record/authority support under time pressure.
Build	Record or synthesize fifteen neutral bench questions in a non-identifying voice. Create factual and doctrinal branches without naming or imitating a real judge.
Student work	Practice ninety-second responses, then annotate a transcript to mark direct answers, concessions, pivots, and authority.
Debrief	Students compare two responses and revise the first sentence of each answer. Faculty provides doctrinal correction.
Guardrails	No automated grade or emotion analysis. Do not equate a particular accent, pitch, or style with professional competence.

More law-school teaching designs

4. Multilingual legal-information audit

Component	Design
Objective	Identify where translation changes jurisdiction, scope, modality, or a legal term of art.
Build	Dub a short, faculty-authored public explainer into a target language. Preserve the source script and machine translation.
Student work	Compare source, translation, and audio; flag terms that require functional rather than literal translation; propose revisions with authority.

Component	Design
Review	A fluent legal expert approves the final version and documents unresolved choices.
Guardrails	Label the dub synthetic; do not use it as individualized legal advice; do not publish an unreviewed translation.

5. Reading roadmap for a flipped class

Component	Design
Objective	Arrive able to identify posture, issue, rule language, reasoning moves, and one ambiguity in the assigned authority.
Build	Produce a two-minute orientation that asks questions but does not provide the case brief. Cite the assigned source in the transcript.
Student work	Submit a low-stakes pre-class note with one passage and one question; use class time for comparison and application.
Feedback	Faculty samples common confusions and revises the next roadmap; the audio is not used to surveil or grade listening behavior.
Guardrails	Transcript is simultaneous; no casebook text is uploaded without rights; students may use the text-only path.

6. Administrative record or legislative-hearing listening lab

Component	Design
Objective	Separate record evidence, policy judgment, statutory interpretation, and rhetorical framing.
Build	Use public-domain or licensed source excerpts, or faculty-authored fictional testimony. Add neutral chapter labels and a reviewed transcript.
Student work	Create a source-indexed claim map and identify what the decisionmaker could lawfully rely on.
Debrief	Return to the original source, not the synthetic narration, to verify quotations and context.
Guardrails	Do not change a real speaker's voice or tone in a way that alters meaning; cite and preserve the original recording.

7. Public scholarship: an accessible audio abstract

Create a three-minute, plain-language summary of published scholarship for a public audience. Link to the paper, name the jurisdiction and limits, include a transcript and sources, disclose synthetic narration, and invite corrections. Do not upload an embargoed manuscript or peer-review material.

Design a student assignment that preserves learning

If students use ElevenLabs, the assignment should assess legal judgment, source use, and reflection—not the realism of a purchased voice. Stanford Teaching Commons recommends beginning with observable

learning objectives, aligning activities and criteria, assessing process, requiring robust citation, and offering guided, low-stakes work.

Ready-to-adapt assignment: synthetic client voicemail

Task

In teams, create a sixty- to ninety-second fictional client voicemail that gives a classmate enough information to identify an initial legal issue but leaves material interview questions unresolved. Use only invented facts and a designed synthetic voice; do not clone any person. Exchange the audio and transcript with another team. The receiving team will produce (1) a fact/assumption chart, (2) eight sequenced follow-up questions, and (3) a two-hundred-word explanation of how tone affected its initial impressions.

Submit

- Final script and transcript
- Audio file with synthetic-media disclosure
- Source/authority list for any legal proposition
- Voice and generation log (tool, model, date, settings)
- One-page process memo identifying changes, errors, and unresolved risks

Alternative

Students may record the faculty-provided script themselves or submit only the written simulation and analysis. The legal-learning criteria are identical.

Assessment criteria

Criterion	Evidence of learning	Suggested weight
Legal issue design	Facts create a plausible issue without assuming the legal conclusion.	25%
Interview strategy	Questions are purposeful, sequenced, non-leading where appropriate, and tied to uncertainty.	30%
Source and accuracy control	Legal propositions and quotations are supported; tool errors are corrected.	20%
Ethics, disclosure, and access	Fictional data, permitted voice, disclosure, accurate transcript, and alternative format are present.	15%
Process reflection	Memo explains choices, revisions, limitations, and what human judgment added.	10%

DO NOT GRADE Do not award points for sounding like a particular gender, race, nationality, age, disability status, region, profession, or social class. Do not infer credibility, effort, emotion, or authorship from synthetic vocal performance.

Stanford teaching guidance: [Integrating AI into assignments](#)

Stanford policy guidance: [Creating your course policy on AI](#)

Faculty learning and productivity uses

For faculty learning

Use	Concrete workflow	What to verify
Hear a draft	Narrate your own public lecture script to detect long sentences, undefined terms, and weak transitions.	Do not treat pleasant delivery as substantive accuracy.
Pronunciation rehearsal	Create short tests for case names, foreign terms, party names, statutes, and acronyms; compare with authoritative pronunciation sources.	The model can confidently mispronounce unfamiliar terms.
Translation audit	Compare a faculty-authored legal-information script with a machine dub and a human-reviewed translation.	Jurisdiction, modality, and terms of art.
Synthetic-media literacy	Create a disclosed fictional clip and document its production trail for evidence or professional-responsibility analysis.	No real-person imitation; secure course distribution.

For productivity with public or Low Risk content

Deliverable	Efficient pattern	Boundary
Course welcome	Faculty-authored sixty-second narration plus transcript and course-policy link.	No student data; update both audio and transcript when policy changes.
Weekly reading roadmap	Reusable script template with source citation and two pre-class questions.	Do not narrate licensed commentary without permission.
Conference or article audio abstract	Public, published content; synthetic voice disclosure; paper link and transcript.	No embargoed or confidential draft.
Public legal education	Plain-language script, expert review, jurisdiction/date limits, citations, transcript, and feedback route.	Not individualized legal advice.
Media rough cut	Studio timeline with placeholder narration and captions for internal review.	Replace or approve before publication; record rights and plan terms.

Uses to keep in Stanford-approved systems

- transcription or summarization of class sessions, faculty meetings, advising, office hours, interviews, or clinics;
- work involving student identity, performance, grades, accommodations, or education records;
- client, privileged, confidential, personnel, human-subject, health, financial, or security information;
- automated legal guidance, student advising, grading, discipline, accommodation decisions, or other high-stakes decisions; and
- voice agents that collect live user data, place calls, or connect to Stanford systems.

CHECK BEFORE EACH TERM Stanford approval and vendor features can change. Recheck the Approved AI Services matrix, Responsible AI guidance, course policy, licensing, and accessibility requirements before reusing this workflow.

Quality-control checklist

Review area	Pass criteria
Purpose	The audio supports a stated learning or communication objective and does not replace the authority students must read.
Data	Only public or authorized Low Risk content was entered; the data-use opt-out is on; no sensitive recording or voice was uploaded.
Rights	You own or have permission for every input, voice, recording, text, image, music reference, and sound.
Identity	No real person is imitated; any clone is your own verified voice; synthetic characters are not deceptive.
Legal accuracy	Names, quotations, citations, holdings, terminology, dates, and jurisdictional limits match authoritative sources.
Audio accuracy	Pronunciation, speaker assignment, emphasis, pace, pauses, and sound cues do not change meaning.
Translation	A fluent legal reviewer approved the target text and listened to the dub.
Disclosure	The audio, player page, and transcript identify AI-generated or altered content.
Accessibility	Audio has an accurate transcript; video has accurate synchronized captions; alternative participation is available.
Documentation	Final script, source list, production log, plan/rights notes, and review date are preserved.
Distribution	Audience and permissions are appropriate; course-only media cannot be mistaken for authentic evidence outside class.

Troubleshooting

Problem	Likely cause	Action
Flat or robotic narration	Voice/model mismatch or overly high stability.	Try a clearer voice, shorten the script, then make a small stability change.
Name or citation is wrong	Unfamiliar token, digits, symbols, or ambiguous spelling.	Write it as spoken; test one sentence; keep standard spelling in the transcript.
Voice changes or adds artifacts	Low stability, style exaggeration, poor source voice, or long generation.	Restore defaults, use style zero, switch voice/model, and generate shorter sections.
Dub sounds fluent but wrong	Literal or context-poor translation.	Edit the target text with a fluent legal reviewer; regenerate only affected segments.
Transcript assigns the wrong speaker	Overlap, noise, or diarization error.	Replay word by word, rename speakers, and compare with the original recording.
Captions differ from final audio	Audio was regenerated after captions were created.	Regenerate or edit captions after the final audio; conduct a full playback review.
Credits or export options differ	Plan or feature changes.	Check current pricing and project settings before production; export and archive promptly.
Students cannot or do not want to use the tool	Cost, privacy, disability, device, or policy barrier.	Use the equivalent text/record-yourself pathway with identical learning criteria.

MAINTENANCE Voice products change quickly. If a screenshot or label no longer matches, update the surrounding instruction—not only the image. Retire references to legacy “Projects,” “Voiceover Studio,” or Dubbing Studio workflows when the current Studio or Dubbing product replaces them.

Official sources and verification notes

Verified August 4, 2026. The tutorial relies primarily on official ElevenLabs product/legal documentation and official Stanford guidance. Links should be rechecked before publication or classroom adoption.

1. [Stanford University IT — Responsible AI at Stanford](#)
2. [Stanford University IT — Approved AI Services](#)
3. [Stanford Office of Digital Accessibility — Audio and Video](#)
4. [Stanford Office of Digital Accessibility — Captioning Basics](#)
5. [Stanford Teaching Commons — Integrating AI into assignments](#)
6. [ElevenLabs — Documentation overview](#)
7. [ElevenLabs — ElevenCreative overview](#)
8. [ElevenLabs — Text to Speech product guide](#)
9. [ElevenLabs — ElevenCreative Studio](#)
10. [ElevenLabs — Dubbing](#)
11. [ElevenLabs — Speech to Text](#)
12. [ElevenLabs — Voice Library](#)
13. [ElevenLabs — Voice Design](#)
14. [ElevenLabs — Voice cloning](#)
15. [ElevenLabs — Sound Effects](#)
16. [ElevenLabs — Music](#)
17. [ElevenLabs — Terms of Service](#)
18. [ElevenLabs — Privacy Policy](#)
19. [ElevenLabs — Prohibited Use Policy](#)
20. [ElevenLabs Help — Data used to improve AI models](#)
21. [ElevenLabs Help — Publishing and plan rights](#)
22. [ElevenLabs — Pricing](#)
23. [ElevenLabs — Audiobooks](#)
24. [ElevenLabs — Image & Video](#)

Limits of this tutorial

This guide does not approve ElevenLabs for Stanford use, interpret vendor terms, provide legal advice, replace a data risk assessment, or authorize processing of protected material. Faculty remain responsible for course policy, academic integrity, accessibility, copyright, privacy, consent, professional obligations, and review of every output.

Appendix A: Sample Text-to-Speech Script

Eleven Laps Text-to-Speech Sample for the Exercise on Page 10.

[measured]

This short audio is a reading roadmap for Tinker versus Des Moines Independent Community School District.

The decision appears in volume three hundred ninety-three of the United States Reports, beginning at page five hundred three. The Supreme Court decided the case in nineteen sixty-nine. Justice Fortas delivered the opinion of the Court.

[short pause]

In December nineteen sixty-five, students in Des Moines, Iowa, planned to wear black armbands. They wanted to express opposition to the war in Vietnam and support for a truce.

School principals learned of the plan. They adopted a policy requiring students to remove the armbands. A student who refused would be suspended.

Mary Beth Tinker, John Tinker, and Christopher Eckhardt wore the armbands. School officials sent them home.

[short pause]

The students sued under Section nineteen eighty-three of Title forty-two of the United States Code. They requested an injunction and nominal damages.

The federal district court dismissed the case. The United States Court of Appeals for the Eighth Circuit divided evenly. That left the dismissal in place.

The Supreme Court then agreed to hear the case.

[long pause]

[curious]

As you read, listen for three questions.

First: Why did the Court treat the armbands as expression closely related to speech?

[short pause]

Second: What evidence would have allowed school officials to predict a material and substantial disruption?

[short pause]

Third: Which statements were necessary to decide this dispute? Which broader statements about student rights might be dicta?

[long pause]

[thoughtful]

Do not stop at the famous statement about students' rights at the schoolhouse gate.

Identify the language that does the analytical work.

The Court explained that schools may regulate student conduct that materially disrupts classwork, produces substantial disorder, or invades the rights of others.

[short pause]

Now test that standard against the record.

The majority found no disruption of classwork. A few students made hostile remarks outside the classrooms. But there were no threats or acts of violence on school grounds.

The record also did not show facts that would have allowed school officials reasonably to forecast substantial disruption or material interference with school activities.

[long pause]

[measured]

The Supreme Court held that, on this record, the school could not prohibit the students' silent and passive expression.

The Court reversed the lower court's judgment. It returned the case for further proceedings.

The Court did not decide what final relief the students should receive.

[short pause]

Bring two items to class.

Bring one passage that supports your explanation of the Court's rule.

And bring one uncertainty about how that rule should apply in a different school setting.

[long pause]

This audio uses an artificial intelligence generated voice.

The script was written and reviewed by [name].

Consult the assigned opinion and course materials for the controlling text.

References (for script)

Tinker v. Des Moines Independent Community School District, 393 U.S. 503 (1969).

<https://www.govinfo.gov/content/pkg/USREPORTS-393/pdf/USREPORTS-393-503.pdf>

42 U.S.C. § 1983 (2018).

<https://uscode.house.gov/view.xhtml?req=granuleid:USC-prelim-title42-section1983>


Appendix B: Video Recordings

Video Recordings:

- [Elven Labs Studio Demo](#)

Eleven Labs Studio Demo

lab studio to create a short



- [Eleven Labs Text-to-Speech Demo](#)

